



MONASH University
Information Technology

FIT5201

Data Analysis Algorithms

Neural Networks

Outline

- Regularization
- Unsupervised and Self-taught Learning

The power of neural networks

- The model class corresponding to neural networks can represent almost any function (given some minor conditions) provided the network has a sufficiently large number of hidden units
 - Have been widely studied
 - 9 layer can solve many low-level intelligence task pretty well

The power of neural networks




- Classification problem
 - Approximate the target decision boundary to any required precision
- Regression problem:
 - Approximate the target function to any precision
- Price:
 - Large number of neurons in the hidden layers
 - Large number of parameters
 - Tend to overfit the training data

The power of neural networks

- Methods to prevent overfitting
 - Use a large training data
 - Use regularization methods (i.e., weight decay)
 - Use deep architecture instead of wide and shallow architecture
 - > Given same number of neurons, deep design performs better
 - > Given same performance, deep architecture needs smaller number of neurons
 - Early stopping


Regularization

- Similar to regression, NNs can overfit the data if too many hidden neurons are added.
 - Blue points: 10 data points drawn from the sinusoidal function plus some noise
 - Red line: the model trained on the neural network with 1 hidden layer
 - > The number of hidden units is different: 1, 3, and 10



Regularization...

- The generalization error, however, is not a simple function of M due to the presence of local minima in the error function



1. For each value of M , try 30 random initializations, how?
2. Polynomial data set, sum-of-square validation error
3. Which one is the best M value? How to judge?
4. One way to choose M in practice

Methods to regularize

- Use different initial weights and find the one that has the smallest validation set error
- Weight decay
- Early stopping

Weight Decay


- Similar to polynomial fitting in regression
- Use large number for M and control the complexity by using a regularization term to the error function
- The effective model complexity is then determined by the choice of the regularization coefficient λ

Early Stopping


- Stop the training when the **network is not learning anymore**

Early Stopping

- Stop the training when the network is not learning anymore
- Stop at the point where the validation set error starts to increase
 - Indicates overfitting



Training error




Test error

Prechelt, Lutz; Genevieve B. Orr (2012-01-01). "Early Stopping - But When?". In Gregoire Montavon, Klaus-Robert Muller (eds.) Neural Networks: Tricks of the Trade. Lecture Notes in Computer Science. Springer Berlin Heidelberg. Pp. 53-67. ISBN 978-3-642-35289-8. Retrieved 2013-12-15

Early Stopping ...

- Has the same effect as weight decay
 - Has theoretical support
 - Intuitive explanations




Unsupervised Learning with NN

- Neural networks are generally considered supervised learning
- In unsupervised learning no model output is present
- Important component in deep learning
- Unsupervised learning:
 - Clustering
 - Dimension reduction
 - Encoding
 -


Unsupervised Learning with NN

- Connect the output back to the input and create a loop
 - Number of output neurons is the same as the number of input neurons
 - Optimize the weights so that the input pattern is also at the output
 - Creates an unsupervised neural network



Autoencoder

- Unsupervised NN can be considered as dimensionality reduction




Minimize
$$E(\theta) := \frac{1}{2} \sum_{n=1}^N \|h_{\theta}(\mathbf{x}_n) - \mathbf{x}_n\|_2^2$$

PCA vs Autoencoder

- If the transformations in the NN is linear, Autoencoder is identical to PCA.
 - They both minimize linear models using the same sum of squares method
 - For linear models, PCA is more efficient than Autoencoder
- Auto encoder can do non-linear transformations but PCA cannot.
 - One hidden layer solution can be given by the projection onto the principal component subspace
 - Need extra hidden layer

Autoencoder ...

- F_1 is the encoder and F_2 is the decoder



Visualization of Autoencoder

- Visualize what one hidden unit encodes
- How?
- What's visualizable?
 - The input
- For each input, one hidden unit will generate one value
- Try all the inputs to see which one can maximize this hidden unit.
 - Expensive to do this
- Solutions:
 - Without running on all the inputs, can we just compute with equations?

Visualization of Autoencoder

$$a_i^{(2)} = f\left(\sum_{j=1}^{100} W_{ij}^{(1)} x_j + b_i^{(1)}\right)$$


1. Treat x_j as variable, W and b as constants
2. Maximize this function

$$x_j = \frac{W_{ij}^{(1)}}{\sqrt{\sum_{j=1}^{100} (W_{ij}^{(1)})^2}}$$

with the constraint: $\|\mathbf{x}\| \leq 1$

Visualization of Autoencoder

- An autoencoder on images
- Visualize the x that maximizes each hidden unit.




Self-Taught Learning with NNs

- If you have limited labels
- Semi-supervised learning
- How to make use of unlabelled data


Self-Taught Learning with NNs

- Learn the features from the autoencoder



Self-Taught Learning with NNs

- Trained encoder can now be used in training a classifier



1. Use a to replace the original input x
2. Concatenate a and x , and use the concatenation to replace x

Want to know more about NNs

