



MONASH
University

MONASH
BUSINESS
SCHOOL

Statistical Thinking (ETC2420/ETC5242)

Associate Professor Catherine Forbes

Week 2: Introduction to data

Week 1 Learning Goals

The learning goals for Week 1:

- Learn how to set up **R** and **RStudio** on your own device.
- Learn to install and load **R** packages.
- Learn what are **R Markdown** files and reproducible research.
- Learn what is 'the tidyverse'.
- Learn some basic **R** commands to manipulate and plot data.

Week 2 Learning Goals

The learning goals for Week 2:

- Identify types of variables, summarise them appropriately, and characterise relationships between them.
- Describe scientific data collection principles.
- Classify variables as being numerical or categorical.
- Illustrate 'tidy data' organisational principles.
- Produce descriptive summaries of numerical and categorical data using appropriate ggplot2, tidyr and dplyr functions.

Assigned reading for Week 2:

- Chapter 1 in ISRS (prescribed textbook: *Introductory Statistics with Randomization and Simulation*)
- Chapters 1, 3 and 12 in *R for Data Science*, especially:
 - ▶ Sections 1.1 - 1.2
 - ▶ Sections 3.1 - 3.4
 - ▶ Section 3.6
 - ▶ Sections 12.1 - 12.3

What is statistical thinking?

- A way of understanding a complex world. . .
 - ▶ using simple terms for essential structure,
 - ▶ acknowledging and assessing degree of uncertainty,
 - ▶ based on foundations from maths, stats, computer science, psychology, and more
- Using data to challenge intuition

"Perhaps H. G. Wells was right when he said 'statistical thinking will one day be as necessary for efficient citizenship as the ability to read and write'!"

- Samuel S Wilks, President of the American Statistical Association in 1951

What can statistics do for us?

Three major things:

- **Describe:**

Characterise complex and noisy data using simpler terms

- **Decide:**

When uncertain, use data to support decisions

- **Predict:**

Anticipate relative chance for potential outcomes of a future random event, based on past data

We'll consider all of these from the perspective of:

- **Frequentist thinking** based on randomisation and simulation

- **Bayesian thinking** using subjective probabilities

But first, let's just focus on the **data**

Upcoming video book chapters

- First 3 learning goals covered by videos produced by the ISRS authors
 - ▶ short and clear
 - ▶ follow the textbook sections in Chapter 1
 - ▶ correspond to Week 2 video book chapters 2 through 8
- Summary points provided on upcoming slides
- After these, we will pick up again with a focus on the final two learning goals

1.1 Case study - using stents to prevent strokes

- Stents and risk of stroke
- Treatment and control groups
- A data table
- Summary statistics
- Random fluctuation

1.2: Data basics

- Observations, variables and data matrices
- Types of variables
- Relationships between variables

1.3 Data collection principles

- Populations and samples
- Sampling from a population
- Explanatory and response variables
- Observational studies and experiments

1.4 Observational studies and sampling strategies

- Observational studies
- Simple random sampling
- Stratified sampling

1.5: Experiments

- Control groups and treatment groups
- Randomisation and replication
- Blocking
- Blinding and placebos

1.6 Examining numerical data

- Scatterplots for paired data
- Dot plots and the mean
- Histograms and shape
- Variance and standard deviation
- Box plots, the median, and robust statistics

1.7 Considering categorical data

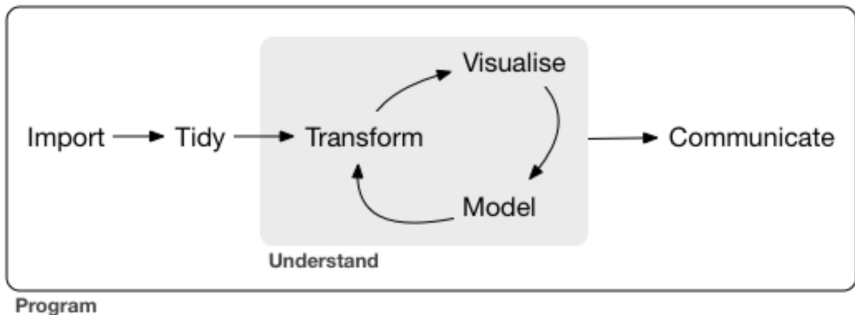
- Contingency tables and bar plots
- Row and column proportions
- Comparing numerical data across groups

Introduction to tidy data

- Recall the **tidyverse**: An **R** package, which is itself comprised of many other **R** packages
 - ▶ ggplot2: data visualisation
 - ▶ dplyr: data manipulation
 - ▶ tidyr: data organisation
 - ▶ readr: data import
 - ▶ purrr: function iteration
 - ▶ tibble: data storage
 - ▶ stringr: string management
 - ▶ forcats: categorical data functions

Why tidy data?

- “A typical data science project” (R4ds)



- Put (rectangular) data in a standard format

Tidy format

- Observations in rows
- Variables in columns
- Values in cells

country	year	cases	population
Afghanistan	1999	745	19987071
Afghanistan	2000	2666	20095360
Brazil	1999	37737	172006362
Brazil	2000	80488	174004898
China	1999	212258	1270015272
China	2000	216766	128002583

variables

country	year	cases	population
Afghanistan	1999	745	19987071
Afghanistan	2000	2666	20095360
Brazil	1999	37737	172006362
Brazil	2000	80488	174004898
China	1999	212258	1270015272
China	2000	216766	128002583

observations

country	year	cases	population
Afghanistan	1999	745	19987071
Afghanistan	2000	2666	20095360
Brazil	1999	37737	172006362
Brazil	2000	80488	174004898
China	1999	212258	1270015272
China	2000	216766	128002583

values

Which data set is in a 'tidy' format?

- All data sets contain WHO counts of TB cases during 1999 and 2000 with population size in 3 countries
- Which data set is in a 'tidy' format?

Same data set four ways (table1)

```
# A tibble: 6 x 4
```

	country	year	cases	population
	<chr>	<int>	<int>	<int>
1	Afghanistan	1999	745	19987071
2	Afghanistan	2000	2666	20595360
3	Brazil	1999	37737	172006362
4	Brazil	2000	80488	174504898
5	China	1999	212258	1272915272
6	China	2000	213766	1280428583

Which data set is in a 'tidy' format?

Same data set four ways (table2)

```
# A tibble: 12 x 4
```

	country	year	type	count
	<chr>	<int>	<chr>	<int>
1	Afghanistan	1999	cases	745
2	Afghanistan	1999	population	19987071
3	Afghanistan	2000	cases	2666
4	Afghanistan	2000	population	20595360
5	Brazil	1999	cases	37737
6	Brazil	1999	population	172006362
7	Brazil	2000	cases	80488
8	Brazil	2000	population	174504898
9	China	1999	cases	212258
10	China	1999	population	1272915272
11	China	2000	cases	213766
12	China	2000	population	1280428583

Which data set is in a 'tidy' format?

Same data set four ways (table3)

```
# A tibble: 6 x 3
  country      year rate
* <chr>      <int> <chr>
1 Afghanistan  1999 745/19987071
2 Afghanistan  2000 2666/20595360
3 Brazil       1999 37737/172006362
4 Brazil       2000 80488/174504898
5 China        1999 212258/1272915272
6 China        2000 213766/1280428583
```

Which data set is in a 'tidy' format?

Table4a

```
# A tibble: 3 x 3
  country      '1999' '2000'
* <chr>      <int>  <int>
1 Afghanistan    745    2666
2 Brazil        37737   80488
3 China         212258  213766
```

Table 4b

```
# A tibble: 3 x 3
  country      '1999'      '2000'
* <chr>      <int>      <int>
1 Afghanistan 19987071  20595360
2 Brazil      172006362  174504898
3 China       1272915272 1280428583
```

Advantages of tidy data

Main package is **tidyr**. Advantages of tidy data include:

1 Consistent workflow

- tools have an underlying uniformity

2 Computational benefits

- mathematical operations on vectors (variables in columns) are fast



3 All the other tidyverse packages work with tidy data

- tibble
- ggplot2
- dplyr

- Making beautiful plots is “easy” with ggplot2
 - ▶ you have already used ggplot2 in Lab 1
- Create data visualisations (plots) based on “The Grammar of Graphics”
- You provide
 - ▶ the tibble containing the data
 - ▶ the mapping of variables to aesthetics
 - ▶ desired geom(s), to define the type of plot
 - ▶ additional layers: stats, scales, coordinate systems, faceting, position adjustments, labels, and legends
- Produces an object that can be stored, with
 - ▶ layers added later
 - ▶ printing executed later
 - ▶ used as input into other function

A simple ggplot2 example


```
library(openintro)
library(gridExtra)
data(arbuthnot)
p1 <- ggplot(data = arbuthnot, aes(x = year, y = girls)) + geom_point()
p2 <- ggplot(data = arbuthnot, aes(x = year, y = girls)) + geom_line()
grid.arrange(p1, p2, ncol = 2)
```



A simple ggplot2 example

- Compare *colour*="blue" in the geom:
- See *mpg* in **R** help for more information


```
ggplot(mpg, aes(displ, hwy)) + geom_point(colour = "blue")
```



A simple ggplot2 example

- With `colour=class` in the aesthetic
 - ▶ legend shows automatically
 - ▶ notice the *pipe* (`%>%`)


```
mpg %>% ggplot(aes(displ, hwy, colour = class)) + geom_point()
```



A simple ggplot2 example

- Change *theme* and size of y-axis label font
- See more options in the R Graphics Cookbook

```
ggplot(mpg, aes(displ, hwy, colour = class)) + geom_point() +  
  xlab("Engine displacement (litres)") + ylab("Highway (miles per gallon)") +  
  theme_bw() + theme(axis.title.y = element_text(size = 8))
```



Data wrangling

- Getting messy data into a standard format is known as **wrangling**
- Having a standard format keeps new variable definitions consistent
- Three main parts to data wrangling:



To **Import**

- Use **readr** package functions, e.g. `read_csv()`
- Use **readxl** package (not in **tidyverse**), e.g. `read_excel()`

Many other packages and functions are useful for data that is really messy

- See Chapters 9-16 in **R** for Data Science for many important tips

Data wrangling

To **Tidy**

- Put data into a *tibble*
- Use **tidyr** functions to reshape *tibble*
 - ▶ *pivot_longer()* can stack columns
 - ▶ *pivot_wider()* can unstack columns
 - ▶ refer to **vignette("tidy-data")** and **vignette("pivot")** to learn more
- Use **dplyr** “verb” functions to manipulate data in your tibble
 - ▶ *filter()*, *slice()*, *arrange()*, *desc()* work on rows
 - ▶ *select()*, *rename()*, *mutate()*, *relocate()* work on columns
 - ▶ *summarise()* collapses a group into a single row
 - ▶ refer to **vignette("dplyr")** to learn more
- The “pipe” operator (`%>%`) from the **magrittr** package is also useful

Vignettes

Most good **R** packages come with one or more **vignette**

- A tutorial to help potential users learn how to use a package
- Find list of vignettes for a package from package site:
 - ▶ <https://cran.r-project.org/web/packages/>
 - ▶ sort by package name
 - ▶ sort by package date
- Consider
 - ▶ <https://cran.r-project.org/web/packages/tidyverse/>
 - ▶ <https://cran.r-project.org/web/packages/tidyr/>
 - ▶ <https://cran.r-project.org/web/packages/dplyr/>

Vignettes will typically include

- An introduction, explaining motivation, or a rational
- Easy-to-replicate examples
- Details of individual functions and available options
- Alerts to conflicts with functions from other packages

RStudio produces **Cheatsheets** for some frequently used activities.

- See <https://rstudio.com/resources/cheatsheets/>
- Also available under the **Help** menu in **RStudio**

You may find these cheatsheets particularly helpful

- Data visualisation (ggplot2)
- Data transformation (dplyr)
- **R Markdown**
- **RStudio**