



墨学教育
—MELBSTUDY—

FIT5047 Week 1

Agent, Search

授课老师: Joe



FIT5047平时班 – Week 1

AGENTS



FIT5047平时班 – Week 1

An **agent** is anything that can be viewed as **perceiving its environment through sensors**, and **acting upon that environment through actuators**

Human agent:

- eyes, ears and other organs for sensors
- hands, legs, mouth and other body parts for actuators

Robotic agent:

- cameras and infrared range finders for sensors
- various motors for actuators



FIT5047平时班 – Week 1

Rationality

*For each possible **percept sequence**, a rational agent should select an **action** that is expected to maximize its **performance measure**, given the evidence provided by the **percept sequence** and the agent's **built-in knowledge***



FIT5047平时班 – Week 1

Task environment - PEAS

- **PEAS**

- Performance measure
- Environment
- Actuators
- Sensors



FIT5047平时班 – Week 1

Automated taxi driver:

- **Performance measure**

- Safe, fast, legal, comfortable trip, minimize fuel consumption, maximize profit

- **Environment**

- Road types, road contents, customers, operating conditions

- **Actuators**

- Control over the car, interfaces for informing other vehicles and informing passengers

- **Sensors**

- Cameras, sonar, speedometer, GPS, odometer, engine sensors, interface for receiving information from other vehicles and passengers (e.g., speech recognizer)



FIT5047平时班 – Week 1

Internet shopping agent:

- **Performance measure**
 - cheap, good quality, appropriate product
- **Environment**
 - current WWW sites, vendors
- **Actuators**
 - display to user, follow URL, fill in form
- **Sensors**
 - HTML pages (text, graphics, scripts)



Environment Attributes

- **Fully (partially) observable** – An agent's sensors give it access to the complete state of the environment at all times
- **Known (unknown)** – An agent knows the “laws” of the environment
- **Single (multi) agent** – An agent operating by itself in an environment
- **Deterministic (stochastic)** – The next state is completely determined by the current state and the action executed by the agent
- **Episodic (sequential)** – The agent's experience is divided into atomic *episodes*. The next episode does NOT depend on previous actions
 - In each episode an agent perceives a percept and performs a single action
- **Static (dynamic)** – The environment is unchanged while an agent is deliberating
- **Discrete (continuous)** – Pertains to number of states, the way time is handled, and number of percepts and actions
 - E.g., state may be continuous, but actions may be discrete



FIT5047平时班 – Week 1

| Agent Type | Action selected based on |
|-----------------|--|
| Simple reflex | current percept |
| Model based | + internal state (world model) |
| Goal based | + goal |
| Utility based | + utility function |
| Learning | performance element = above agent + critic + learning element + problem generator (exploratory) |



FIT5047平时班 – Week 1

| Agent Type | Action |
|--|---|
| Simple reflex | brake when brake-lights of car in front light up |
| Model based | + remember the roads travelled, time, state |
| Goal based | + make a plan to reach a destination |
| Utility based | + quickest with least petrol consumption |
| Learning performance elem + critic + learning element + problem generator | above agent observes the world & informs learning elem formulates new driving rules based on the feedback from the critic might suggest some driving exercises |



Search



Problem formulation

- Initial state
- Actions
- Constraints
- Goal test
- Path cost function



FIT5047平时班 – Week 1

1. initial state, e.g., “at Arad”

2. actions

- e.g., {Go(Sibiu), Go(Timisoara), ... }

transition model

- e.g., $Result(In(Arad), Go(Zerind)) \rightarrow In(Zerind)$

3. constraints – nil

4. goal test can be


- explicit, e.g., $In(Bucharest)$
- implicit, e.g., $Checkmate(x)$

5. path cost (additive)

- e.g., sum of distances, number of actions executed
- $c(s, a, s')$ is the step cost of taking action a at state s to reach state s' (assumed to be ≥ 0)



FIT5047平时班 – Week 1





FIT5047平时班 – Week 1

- **States**
 - Location of each of the 8 tiles in one of the 9 squares
- **Operators**
 - Possible moves of blank tile
- **Constraints**
 - A tile cannot move out of bounds
- **Goal test**
 - Have we reached the goal configuration?
- **Path cost**
 - If we want to minimize the number of steps, then cost of 1 per step



Searching strategies

Backtracking – at any point in time, we keep one path only

- If we fail, we go back to the last decision point and **erase the failed path**
- Backtracking occurs when
 1. we reach a DEADEND state OR
 2. there are no more applicable rules OR
 3. we generate a previously encountered state description OR
 4. an arbitrary number of rules has been applied without reaching the goal



Searching strategies

Backtracking – at any point in time, we keep one path only


- If we fail, we go back to the last decision point and **erase the failed path**
- Backtracking occurs when
 1. we reach a DEADEND state OR
 2. there are no more applicable rules OR
 3. we generate a previously encountered state description OR
 4. an arbitrary number of rules has been applied without reaching the goal

Graphsearch – we keep track of several paths simultaneously

- Done using a structure called a ***search tree/graph***




FIT5047平时班 – Week 1





FIT5047平时班 – Week 1

4 Queens Problem





FIT5047平时班 – Week 1

Tree Search

- Initialize the frontier using the initial state of *problem*
- **Loop**
 1. if the frontier is empty **then return** failure
 2. **choose** a leaf node and remove it from the frontier
 3. if the node contains a goal state **then return** the corresponding solution
 4. **expand** the chosen node, **adding** the resulting nodes to the frontier
- end**


Graph Search

- Initialize the frontier using the initial state of *problem*
- Initialize the *explored set (closed)* to empty
- **Loop**
 1. if the frontier is empty **then return** failure
 2. **choose** a leaf node and remove it from the frontier
 3. if the node contains a goal state **then return** the corresponding solution
 4. add the node to the *explored set*
 5. **expand** the chosen node, **merging** the resulting nodes with the frontier *or the explored set*




FIT5047平时班 – Week 1

Breadth first Search



Depth first search





FIT5047平时班 – Week 1


Uniform Cost Search

1. **if** the frontier is empty **then return** failure
2. **choose** the lowest-cost node in the frontier and remove it from the frontier
3. **if** the node contains a goal state **then return** the corresponding solution
4. **expand** the chosen node
 - a. **if** the resulting nodes are not in the frontier **then add** them to the frontier
 - b. **else if** the resulting nodes are in the frontier **with higher path cost then replace** them with the new nodes



FIT5047平时班 – Week 1

Uniform Cost Search





FIT5047平时班 – Week 1

Depth limited search = Depth first search with limited depth




FIT5047平时班 – Week 1

Iterative deepening search


Limit = 1



Limit = 2



Limit = 3





FIT5047平时班 – Week 1

Best First Search

- Heuristic function (prediction of how good that state is), denoted by h
- Do not care about the actual cost incurred, only looks at the heuristic when making decisions



FIT5047平时班 – Week 1

Best First Search

- Heuristic function (prediction of how good that state is), denoted by h
- Do not care about the actual cost incurred, only looks at the heuristic when making decisions



Algorithm A

- Graphsearch using the evaluation function
 $f(n) = g(n) + h(n)$
- Expands next the node in the frontier with the smallest value of $f(n)$

$g(n)$ = *Actual cost incurred so far*

$h(n)$ = *prediction to goal*

Algorithm A*

- *A*
- *Admissibility of h*
- *Monotonicity of h*



FIT5047平时班 – Week 1

Algorithm A*

- A
- *Admissibility of h*
- *Monotonicity of h*

Admissibility of h


- *Never overestimated to true cost of going to the goal from that node.*
- *If h is not admissible, then the path found is not guaranteed to be optimal.*

Monotonicity fo h

- *h of parent \leq h of child + cost of going from parent to child*



FIT5047平时班 – Week 1

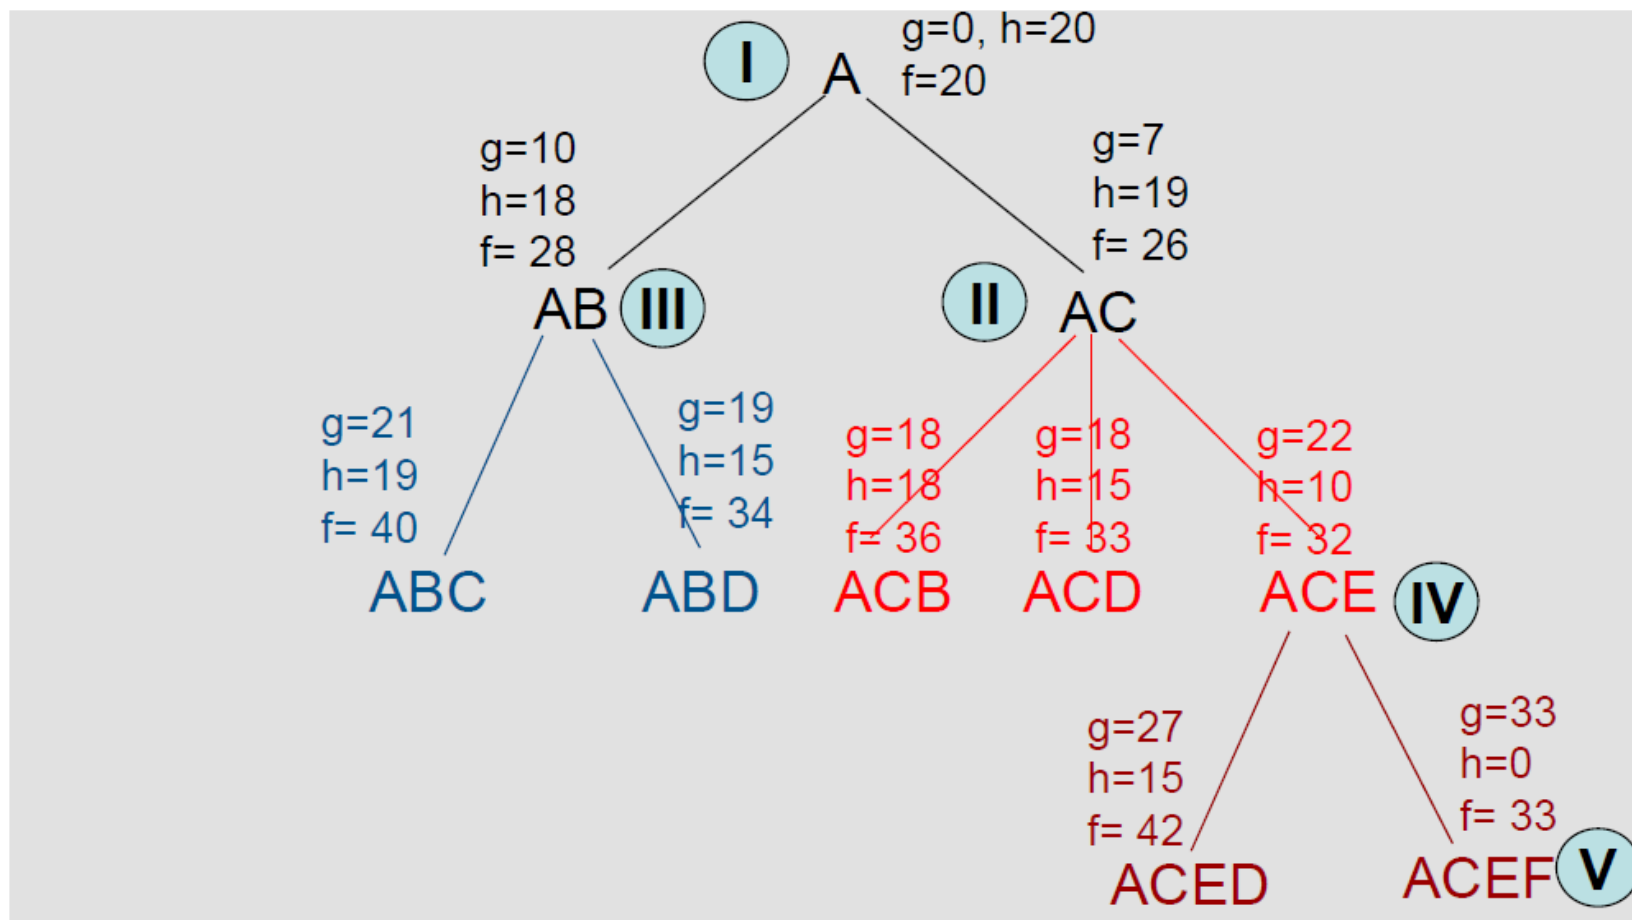


| ROAD DISTANCES | | | | | | |
|----------------|---|----|----|----|----|----|
| | A | B | C | D | E | F |
| A | | 10 | 7 | | | |
| B | | | 11 | 9 | | |
| C | | | | 11 | 15 | |
| D | | | | | 5 | 16 |
| E | | | | | | 11 |

| AIR DISTANCES | | | | | | |
|---------------|---|---|---|---|----|----|
| | A | B | C | D | E | F |
| A | | 4 | 3 | 8 | 12 | 20 |
| B | | | 6 | 5 | 9 | 18 |
| C | | | | 7 | 10 | 19 |
| D | | | | | 5 | 15 |
| E | | | | | | 10 |



FIT5047平时班 – Week 1





Admissible heuristics: 8 Puzzle

- $h_1(n)$ = number of misplaced tiles
- $h_2(n)$ = total *Manhattan distance* (# of squares from desired location of each tile)

- $h_1(S) = ?$
- $h_2(S) = ?$

| | | |
|---|---|---|
| 7 | 2 | 4 |
| 5 | | 6 |
| 8 | 3 | 1 |

Start State

| | | |
|---|---|---|
| | 1 | 2 |
| 3 | 4 | 5 |
| 6 | 7 | 8 |

Goal State



Dominance


- Given two admissible heuristics h_1 and h_2 , if $h_2(n) \geq h_1(n)$ for all n then h_2 dominates h_1
→ h_2 is better for search



FIT5047平时班 – Week 1

Hill Climbing

- Look around, find the next best move , go to that place
- If there is none, then terminate
- Greedy in nature, since it always find something better





FIT5047平时班 – Week 1

Local Beam Search

- Running multiple hill climbing simultaneously



FIT5047平时班 – Week 1

Simulated Annealing

- Similar to hill climbing, but have a chance of going downhill (accept a worse state)
- The probability of accepting a worse state is controlled by **temperature** and **annealing schedule**, which are defined by you
- **$Pr = e^{-\Delta E / T}$**
- When temperature decreases to 0, simulated annealing is just hill climbing



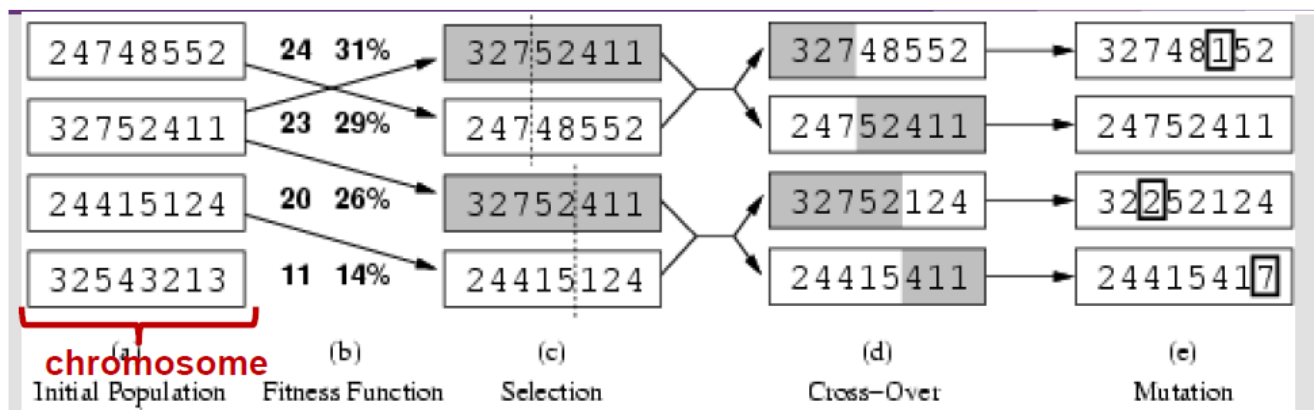
FIT5047平时班 – Week 1

Genetic Algorithm

- The problem is so large and complex, we have no idea where to go
- Start with a population of k randomly generated states (parent)
- A state (chromosome) is represented as a string over a finite alphabet of genes (often a string of 0s and 1s)
- A successor state is generated by combining two parent states
- Evaluation function (fitness function):
 - Higher values for better states
- Produce the next generation of states by selection, crossover and mutation



FIT5047平时班 – Week 1






FIT5047平时班 – Week 1

Adversarial Search

Two players

- Min and Max
- Min wants to minimize the score while Max wants to maximize the score

Minimax ideas





FIT5047平时班 – Week 1

Minimax optimization – Alpha and beta pruning

