MONASH University Information Technology



COMMONWEALTH OF AUSTRALIA

Copyright Regulations 1969

WARNING

This material has been reproduced and communicated to you by or on behalf of Monash University pursuant to Part VB of the Copyright Act 1968 (the Act).

The material in this communication may be subject to copyright under the Act. Any further reproduction or communication of this material by you may be the subject of copyright protection under the Act.

Do not remove this notice.





FIT5047 – Intelligent Systems

Machine Learning
Chapters 18.1-4, 18.6.1,
18.8.1, 20.1-2

Learning Objectives (I)

Machine learning – definition

- Relationship with data mining
- Issues about learning from data

Approaches to learning from data

- Verification driven
- Discovery driven inductive learning
 - > supervised and unsupervised learning



Learning Objectives (II)

Supervised machine learning

- Decision trees
- Naïve Bayes
- k Nearest Neighbour (k-NN)
- Regression
- (Logistic regression)

Clustering (Unsupervised learning)

- The clustering problem
- Similarity measures
- The K-means algorithm



Why Learning?

- Learning is essential for unknown environments
 - e.g., when the designer lacks omniscience
- Learning is necessary in dynamic environments
 - An agent can adapt to changes in the environment not foreseen at design time
- Learning is useful as a system construction method
 - Expose the agent to reality rather than trying to approximate it through equations
- Learning modifies the agent's decision mechanisms to improve performance

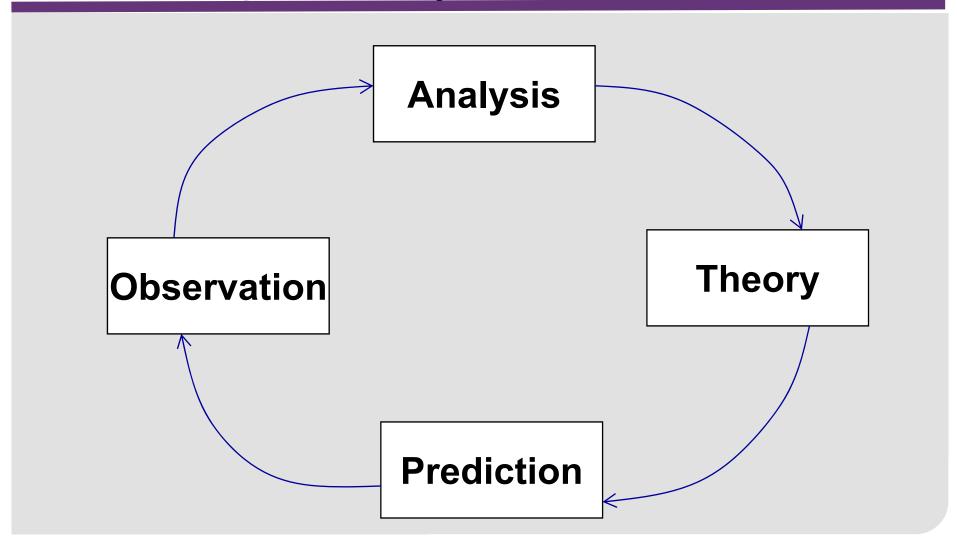


Machine Learning

- Machine Learning is an area of Al concerned with programs that learn from data
 - Machine learning performs inductive learning
- Inductive learning is learning from examples
- A program makes predictions, then learns from feedback on the correctness of the predictions
- Machine Learning is used in many applications
 - E.g., speech recognition, robot training, game playing, and classification of astronomical structures



The Empirical Cycle





Types of Learning Approaches

Verification Driven

 An expert formulates a hypothesis, which is validated by analyzing the data

Discovery Driven – Machine Learning

- Hypotheses are formulated bottom-up from the data
 - Interesting patterns in the data are noticed and evaluated by the user/expert
- Types of discovery-driven learning:
 - > Supervised learning, unsupervised learning, reinforcement learning



Learning Concepts with Machine Learning

- Machine learning aims to reveal knowledge about data
 - concepts, definitions, hypotheses
- A learning algorithm considers many instances and infers useful regularities about them
 - → It can *infer (learn)* a definition of a class
 - > this definition may be viewed as patterns or structures within the data
- The structure of the discovered patterns depends on the learning techniques
 - E.g., rules, decision trees



Completeness and Consistency in Inductive Learning

A definition is

- complete if it recognizes all instances of a class
- consistent if it does not classify any negative examples as instances of this class

An incomplete definition is too narrow

 it would not recognize some "in-class" instances as belonging to the class

An inconsistent definition is too broad

 it would classify some "out-of-class" instances as belonging to the class



Machine Learning as Search

- Find the best hypothesis among many
 - Need a measure of the quality of candidate hypotheses
- The search space is very large
- We don't know the optimal answer



Hypothesis Characteristics

Performance

how often the hypothesis is correct (different measures)

Significance

- must perform better than the naïve hypothesis (e.g., majority)
- statistical significance perform better than chance

Information content

must take into account all relevant information

Transparency

how easy it is to understand the generated hypotheses



Machine Learning as Compression

- We have learnt something if we have created a description of the data that is shorter than the original data set
- Example Input/Output
 - Input: a table of instances
 - Output: a prediction of an instance given a set of features or a suggested clustering of the instances



Machine Learning and Data Mining

- Machine learning (ML) and data mining (DM) overlap
- We can consider ML as a scientific discipline, and DM as an engineering discipline



Machine Learning Principles

• "A general law can never be verified by a finite number of observations. It can, however, be falsified by only one observation."

Karl Popper

- The patterns that machine learning algorithms find can never be definitive theories
- Any results discovered must be tested for performance and statistical significance



Types of Machine Learning

- Supervised learning: correct answers are provided for each input
 - E.g., Decision Trees, Naïve Bayes, K-Nearest
 Neighbour (k-NN), Regression, Neural Nets
- Unsupervised learning: correct answers are not given, must discover patterns in input data
 - E.g., K-Means, Snob (Minimum Message Length Principle)
- (Reinforcement learning: occasional rewards (or punishments) are given
 - E.g., Q-learning







FIT5047 – Intelligent Systems

Mathematical Principles of Machine Learning

Supervised Learning Example: Handwritten Digit Recognition (I)

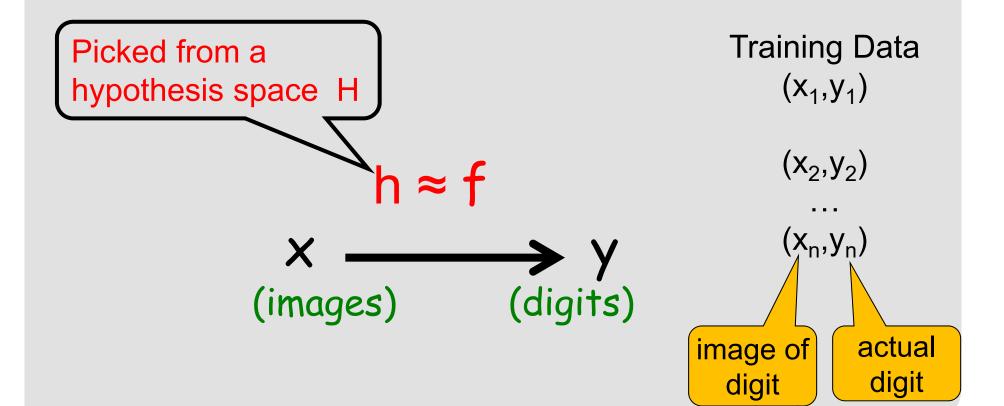
- Classify images into one of the digits: 0, 1, ..., 9
 - Useful for the post office to automatically detect addresses

۵	0	Ø	Ö	Ð	ŏ	Ø	0	८)	Ø	0
((Į.	Ļ	})	Ų	Ų.	ſ)	1
J	2	ð	2.	ఎ	2	\mathcal{Z}	\mathfrak{Q}	Q	2	2
3	3	3	3	3	૭	3	3	3	3	3
Ų,	4	Ų	$oldsymbol{arphi}$	4	4	4	4	4_	4	4
\$	5	5	5	ও	5	5	5	5	5	5
6	6	6	6	6	6	6	G	6	Le .	6
7	7	7	7	7	1	フ	7	7	7	7
8	8	ક	\mathbf{z}	8	8	8	જ	8	8	8
9	9	9	9	9	9	9	9	9	q	9



Supervised Learning Example: Handwritten Digit Recognition (II)

Learn a function f from examples





Supervised Inductive Learning

Setup:

- f is the unknown target function
- We are given some sample pairs from it (x, f(x))
- Problem: learn a function hypothesis h
 - Based on the set of *training* examples
 - Such that h ≈ f (h approximates f as best as possible)
 - > h must generalize well on unseen examples



Picking the Best Hypothesis h

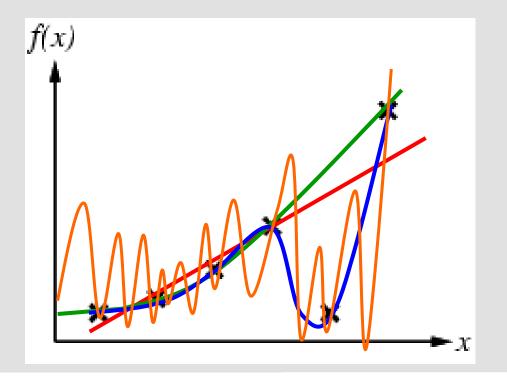
- Big Idea 1:
 Pick h from the space H which agrees with f on the training set
 - Complete and consistent



Inductive Learning – Example (I)

Curve fitting (regression):

h= Straight line? Quadratic? Cubic? Other?



Training Data (x_1,y_1)

 (x_2,y_2)

 (x_n, y_n)



Ockham's Razor

 "Entities must not be multiplied beyond necessity" ("Non sunt multiplicanda entia sine necessitate")

William of Ockham (1287–1347)

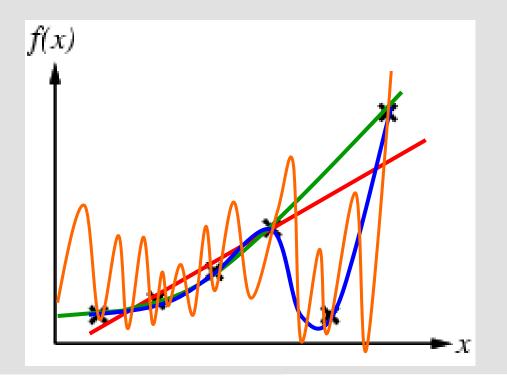
- Big Idea 2: Prefer a simpler hypothesis to complex ones provided both explain the data equally well
 - Overly complex hypotheses tend to overfit the data



Inductive Learning – Example (II)

Curve fitting (regression):

h= Straight line? Quadratic? Cubic? Other?



Training Data (x_1,y_1)

 (x_2, y_2)

 (x_n, y_n)



Mathematical Principles of Learning (I)

Idea 2: To improve generalizability and prevent overfitting



Choose the <u>simplest</u> hypothesis from H which is <u>complete</u> and <u>consistent</u> with the training data

Idea 1: Should be similar to the unknown true underlying function



Bias and Variance

- Bias is the true error of the best classifier in the concept class
 - Bias is high if the concept class cannot model the true data distribution well, e.g., it is too simple
 - High bias → both training and test error are high
- Variance is the error of a trained classifier with respect to the best classifier in the concept class
 - Variance decreases with more training data, and increases with more complicated classifiers
 - High variance → training error is low, and test error is high



Mathematical Principles of Learning (II)

Best learned hypothesis

Idea 2: penalize complexity – regularization

$$h^* = \operatorname{argmin}_{h \in H} \left\{ \left[\sum_{(x,y) \in D} error(h(x), y) \right] \right\}$$

 $+ \lambda Complexity(h)$

Idea 1: Error on the training data

Learning is indeed Search/Optimization







FIT5047 – Intelligent Systems

Concepts in Machine Learning

Types of Data

 Data: (un)labeled instances Used to train the model **Training** Data Used to fine-tune the model **Validation** Data Test Used to measure the Data generalization of the model



Inductive Learning Process

- Features: attribute-value pairs which characterize each instance
- Experimentation cycle
 - Learn model parameters on Training Data
 - Fine tune the model on Validation (Held-out)
 Data
 - Compute performance on Test Data

Very important: never "peek" at the test set!

Learn model (training data) Fine tune (validation data)

Training Data

Validation Data

> Test Data

Evaluate (test data)



Evaluation Metrics

accuracy

correctly predicted predicted

If 80 predictions are correct out of 100 then accuracy = 80/100 = 0.8

- recall
- instances in class C
- precision

correctly predicted as class C
predicted as class C

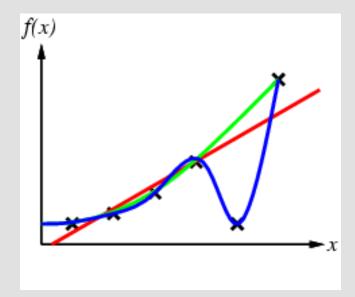
Training Data

Validation Data

> Test Data

Overfitting and Generalization

- We want a learned procedure that does well on test data
 - Overfitting: fitting the training data very closely, but not generalizing well



Training Data

Validation Data

> Test Data







FIT5047 – Intelligent Systems

Approaches to Learning from Data

Discovery Driven – Supervised Learning

- Build patterns by inferring an unknown attribute given the values of known attributes
- Principal Techniques: Classification and Regression
- Supervision:
 - a set of observations (*training data*) is accompanied by labels indicating the class of each observation
 - new data are classified/predicted based on the model learned from the training data

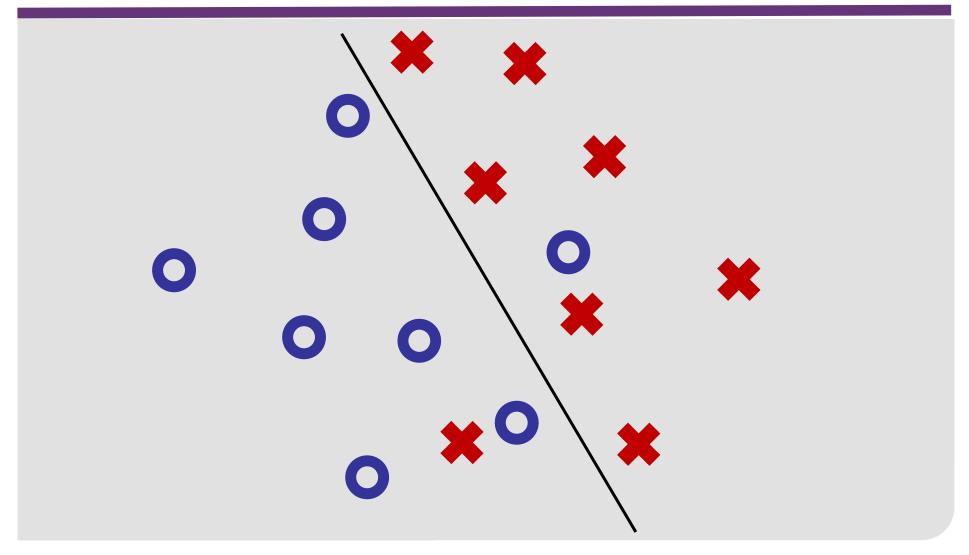


Type of Inferred Value – Classification vs Regression

- Classification
 - Infers categorical or discrete values
- Regression
 - Infers continuous or ordered values



Classification





Classification – Example (I)





















Cat

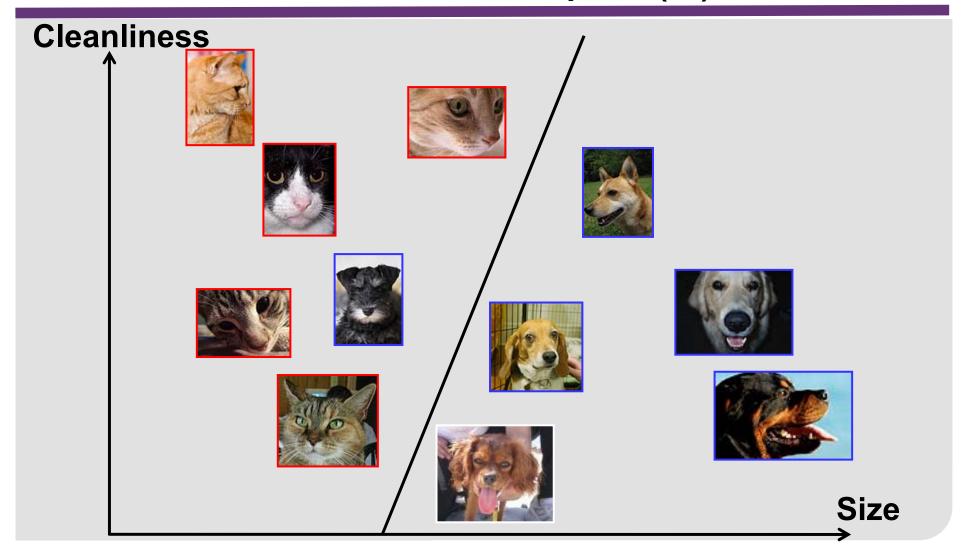




Dog



Classification – Example (II)





Regression – Example (I)

\$

\$\$

\$\$\$

\$\$\$\$









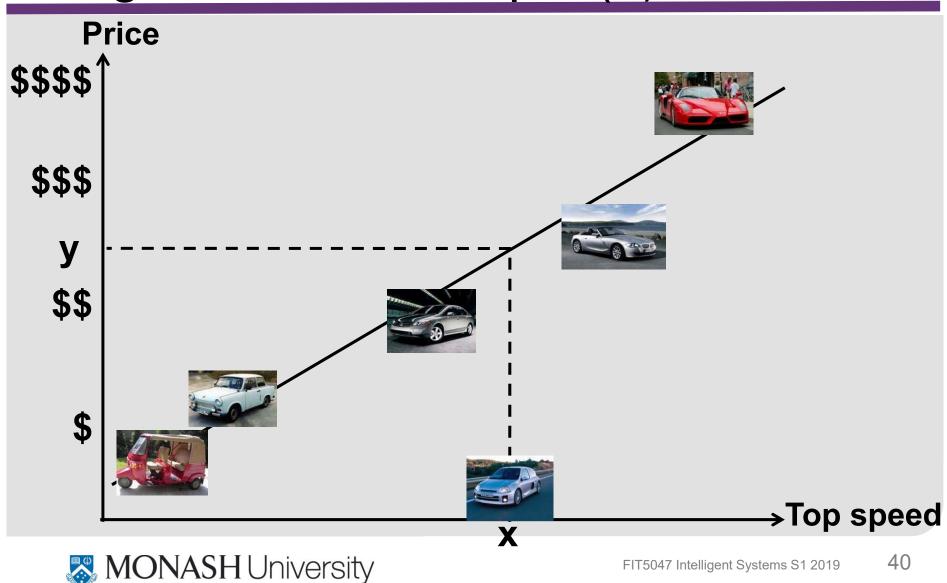








Regression – Example (II)



Time of Inferred Value – Classification vs Prediction

- Classification considers the features of an object in order to assign it to a pre-defined class
 - Examples:
 - > spotting fraudulent insurance claims
 - > identifying high-value customers
- Prediction is the classification of future events
 - Historical data is used to build a model that "explains" outputs for known inputs
 - The model is applied to new inputs to predict future outputs
 - > Examples:
 - predicting which customers won't default on a loan
 - predicting which footy team will win this weekend



Discovery Driven - Unsupervised Learning

- Return "interesting" patterns in the data
- Principal Techniques: Clustering and Association Analysis
- Lack of supervision:
 - given a set of observations (*training data*), infer classes or clusters in the data
 - training data is unlabelled there are no pre-defined classes



Clustering vs Association Analysis

Clustering –

- Groups records of similar items
 - The user must attach meaning to the clusters formed
- Example application
 - > Identify different types of customers

Association analysis –

- Discovers relations hidden in the data
 - > represented in the form of association rules or sets of frequent items
- Example application
 - > Market basket analysis, e.g., diapers -> beer







FIT5047 –Intelligent Systems

Supervised Machine Learning – Classification

Classification Process

1. Model construction

builds a model using a *Training Set* – a set of tuples,
 each comprising a list of features and a class

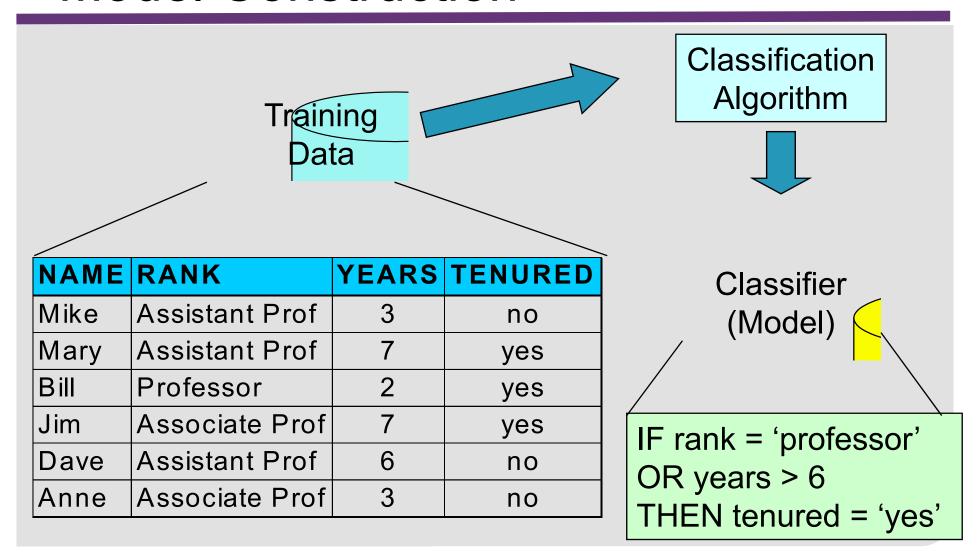
2. Model evaluation – estimates the performance of the model

- uses a *Test Set*, which is different from the training set
- the known label of the items in the test set are compared with the class selected by the model
- Performance can be measured by accuracy, recall, precision

3. Model usage – classifies or predicts the class of unseen items

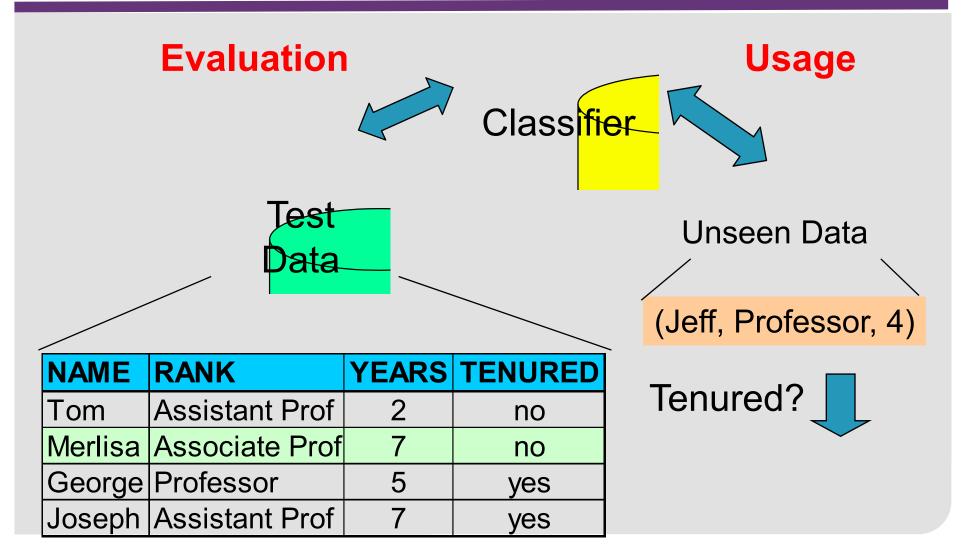


Model Construction





Model Evaluation and Usage





Features of Classification Methods (I)

Hypothesis features

- Performance
 - > Different measures, e.g., accuracy, recall, precision
- Significance
 - > Must perform better than the naïve hypothesis
- Model size (related to information content)
 - > E.g., decision tree size, compactness of classification rules
- Transparency (interpretability)
 - > Ease of understanding the model



Features of Classification Methods (II)

Practicalities of the model

- Robustness
 - > Ability to handle noise and missing values
- Scalability
 - > Speed time to build the model, and time to use the model
 - > Size effect of the size of the dataset on the model



Evaluating Classifiers

Performance

- depends on the representativeness of the training data
- determined using a test set

Need to perform cross-validation, Why?

- X-validation repeated experiments on different test sets
 - > separate the dataset into training and test sets
 - > build the model from the training set, and compute performance on the test set
 - > usually 10-fold or 5-fold X-validation
 - > common set ups: 90/10%; 80/20%; leave-one-out
- Stratified X-validation the test sets are proportional to the classes in the data
 - > e.g., 20% +ive, 80% -ive







FIT5047 –Intelligent Systems

Decision Trees

Decision Trees

- Classify objects based on the values of their explanatory attributes
 - the target classes (dependent attributes) are pre-defined
- Classification is based on a tree structure
 - each non-leaf node is a decision node
 - each leaf node represents a class
- To classify an object
 - each decision node (starting from the root) compares an attribute of the object with a specific attribute value (or range)
 - a path from the root to a leaf node gives the class of the object



Decision Tree Example: Attributes



Input x (vector of values for explanatory attributes)

Output y (target / dependent attribute)



Classification by DT Induction

- Training data records of items that have explanatory attributes and a dependent attribute
 - Input x: Represented by a vector of attribute values
 - Output y: The corresponding target value
- Tree construction algorithm builds a decision tree
 - Top-down, recursive divide-and-conquer approach
 - Based on a greedy algorithm

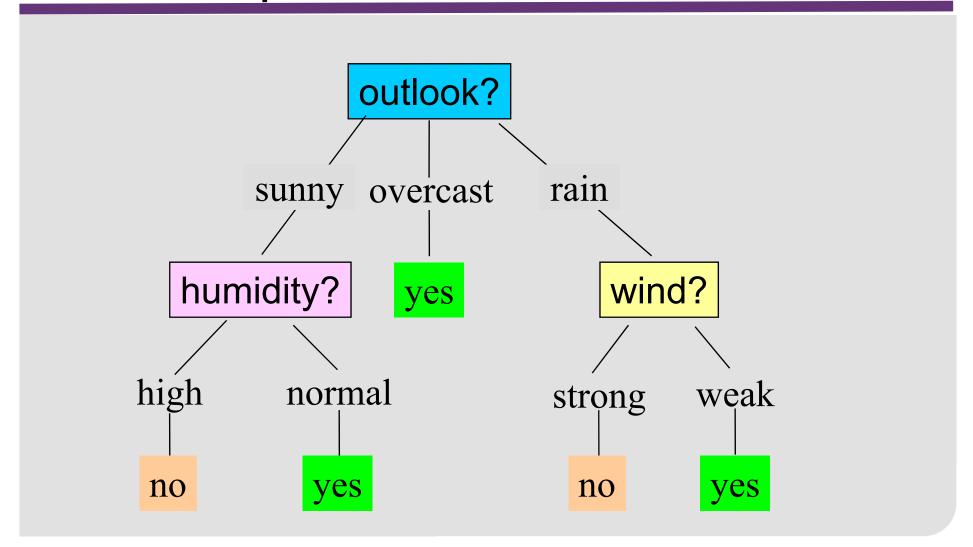


DT Example – Training Dataset

Day	Outlook	Temperature	Humidity	Wind	Play ball
D1	Sunny	Hot	High	Weak	No
D2	Sunny	Hot	High	Strong	No
D3	Overcast	Hot	High	Weak	Yes
D4	Rain	Mild	High	Weak	Yes
D5	Rain	Cool	Normal	Weak	Yes
D6	Rain	Cool	Normal	Strong	No
D7	Overcast	Cool	Normal	Strong	Yes
D8	Sunny	Mild	High	Weak	No
D9	Sunny	Cool	Normal	Weak	Yes
D10	Rain	Mild	Normal	Weak	Yes
D11	Sunny	Mild	Normal	Strong	Yes
D12	Overcast	Mild	High	Strong	Yes
D13	Overcast	Hot	Normal	Weak	Yes
D14	Rain	Mild	High	Strong	No



DT Example – Learned Model





Decision Tree Learning Algorithm (I)

- 1. Start with all training examples at the root
- 2. Partition examples recursively based on selected attributes
 - attributes are categorical
 - > if continuous-valued, they are broken up into ranges
 - attributes are selected using heuristics or a statistical measure
 - > e.g., information gain
- 3. Stop partitioning when there is no further gain in partitioning

Employ majority voting for classifying the leafs



Choosing an Attribute for Splitting

- Approach choose the attribute that best separates training examples into targeted classes
- Examples of proposed techniques
 - Information Gain [Quinlan, 1975]
 - Message Length [Wallace and Patrick, 1993]



What is the "simplest" Tree? – Example

Always predict "yes"

A tree with one node

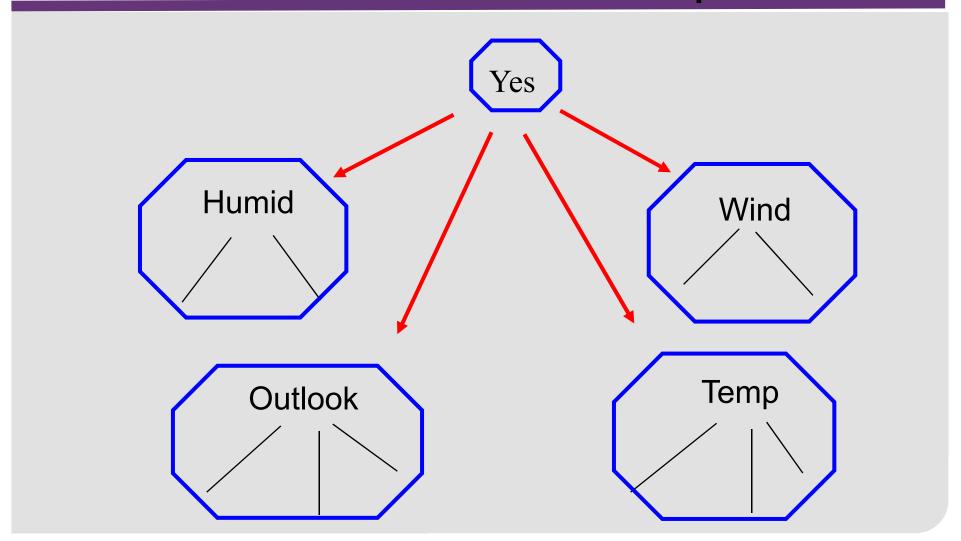
How good it is?

- Correct on 9 examples
- Incorrect on 5 examples
- Notation: [9+,5-]

,	Day	Outlook	Tempe rature	Humi dity	Wind	Play ball
	D1	Sunny	Hot	High	Weak	No
Ī	D2	Sunny	Hot	High	Strong	No
Ī	D3	Overcast	Hot	High	Weak	Yes
	D4	Rain	Mild	High	Weak	Yes
Ī	D5	Rain	Cool	Normal	Weak	Yes
Ī	D6	Rain	Cool	Normal	Strong	No
Ī	D7	Overcast	Cool	Normal	Strong	Yes
Ī	D8	Sunny	Mild	High	Weak	No
Ī	D9	Sunny	Cool	Normal	Weak	Yes
Ī	D10	Rain	Mild	Normal	Weak	Yes
	D11	Sunny	Mild	Normal	Strong	Yes
	D12	Overcast	Mild	High	Strong	Yes
	D13	Overcast	Hot	Normal	Weak	Yes
	D14	Rain	Mild	High	Strong	No

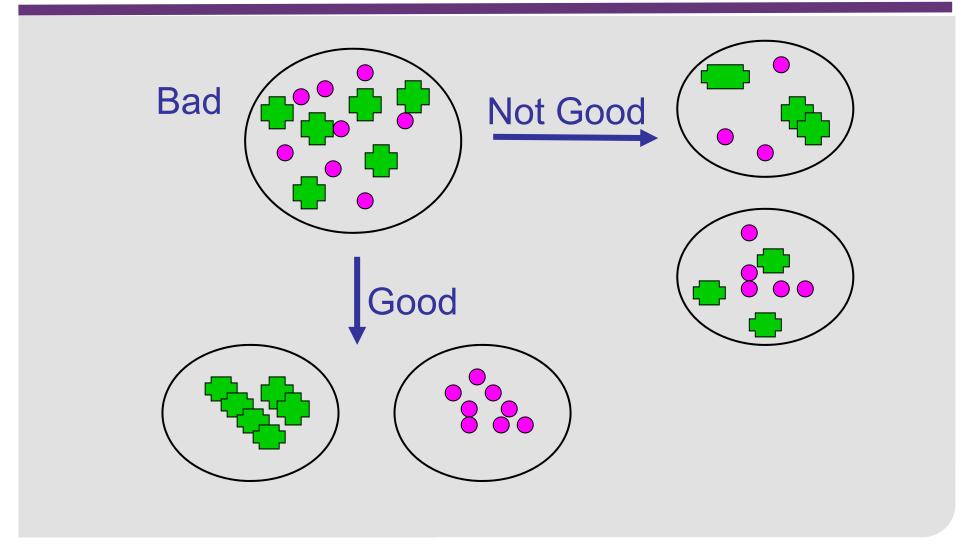


On which Attribute do we Split?





Disorder is Bad, Homogeneity is Good





Using Information Theory to Quantify Uncertainty – Entropy

- Entropy measures the amount of uncertainty in a probability distribution
- Given a discrete random variable on a finite set $X=\{x_1, ..., x_n\}$, with probability distribution function Pr(x)=Pr(X=x), the <u>entropy</u> H(X) of X is defined as

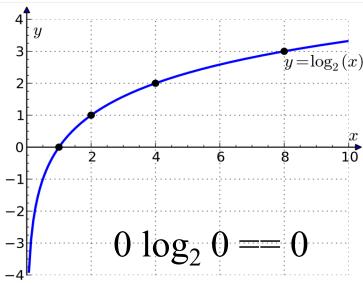
$$H(X) = -\sum_{i=1}^{n} \Pr(x_i) \log_2 \Pr(x_i)$$

where $0 \le H(X)$



Entropy – Example

$$H(X) = -\sum_{i=1}^{n} \Pr(x_i) \log_2 \Pr(x_i)$$

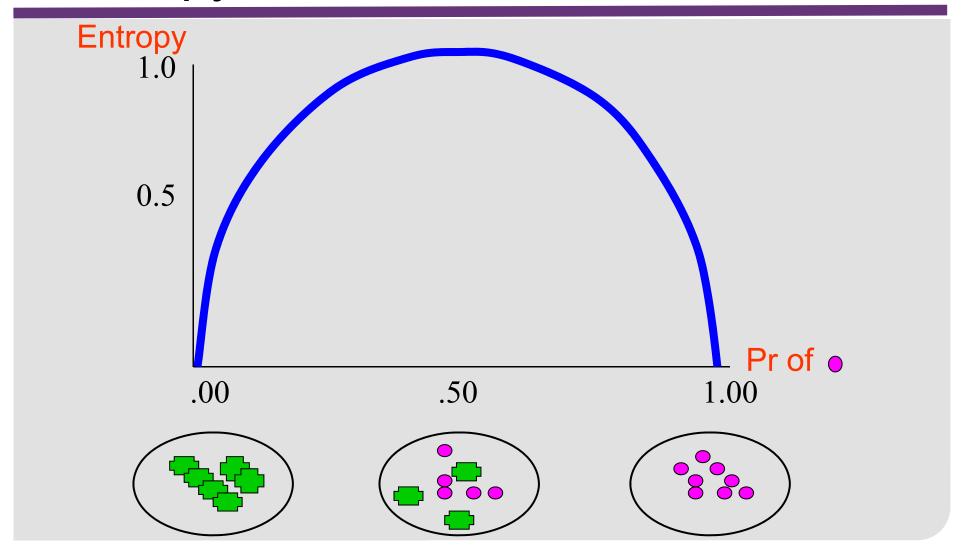


- Suppose there is a random

 variable S that has value a or b
 - Let Pr(a) = 1 and Pr(b) = 0
 - Let Pr(a) = 0.9 and Pr(b) = 0.1
 - Let Pr(a) = 0.5 and Pr(b) = 0.5
 - Which probability assignment <u>maximizes</u> H(S)?



Entropy





Entropy is Bad, Homogeneity is Good

Let S be a set of examples

- Labeled positive or negative
- Entropy(S) = -P $log_2(P)$ N $log_2(N)$
 - P is the proportion of positive examples
 - N is the proportion of negative examples
 - $0 \log 0 == 0$
- Example: S has 9 pos and 5 neg
 - Entropy([9+, 5-])

$$= -\frac{9}{14} \log_2 \frac{9}{14} - \frac{5}{14} \log_2 \frac{5}{14} = 0.94$$



Information Gain (I)

• Information gain is a measure of the expected reduction in entropy resulting from splitting along attribute *A*

$$IG(S,A) = H(S) - \sum_{v \in Values(A)} \frac{S_v}{S} H(S_v)$$

Expected value of H from splitting on A

where

- -v is a value for A (we sum over all values)
- $-S_v$ is a subset of S for which attribute A has value v



Using Information Gain to Learn a DT

- The same attributes must describe each sample
- Attributes are assumed to be categorical (for now)
- Select the attribute with the highest Information
 Gain → largest reduction in entropy



Information Gain of Splitting on Wind

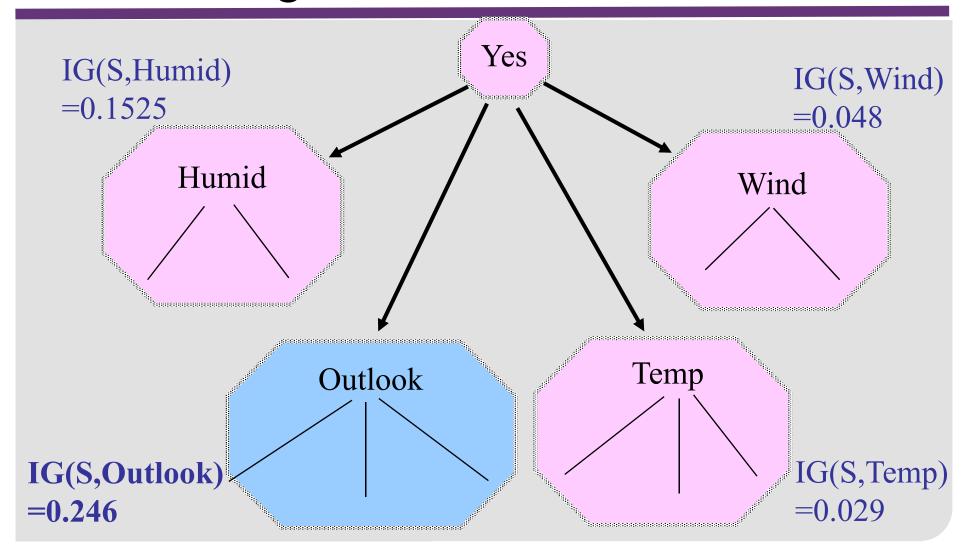
```
S = [9+, 5-]
S_{\text{weak}} = [6+, 2-]
S_{\text{strong}} = [3+, 3-]
    IG(S, wind)
    = H(S) - \sum_{|S|} \frac{|S_v|}{|S|} H(S_v)
                 v∈{weak,strong}
= H(S) - 8/14 H(S_{weak}) - 6/14 H(S_{strong})
= 0.94 - (8/14) 0.811 - (6/14) 1.00
= 0.048
```

Day	Wind F	Play ball?)
d1	weak	no	
d2	strong	no	
d3	weak	yes	
d4	weak	yes	
d5	weak	yes	
d6	strong	no	
d7	strong	yes	
d8	weak	no	
d9	weak	yes	
d10	weak	yes	
d11	strong	yes	
d12	strong	yes	
d13	weak	yes	
d14	strong	no	



Values(wind)=weak, strong

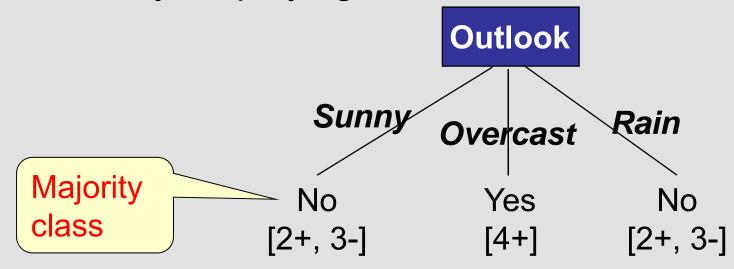
Evaluating Attributes





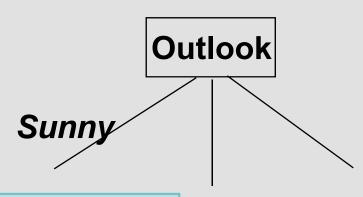
Building a Decision Tree (I)

Good day for playing ball?





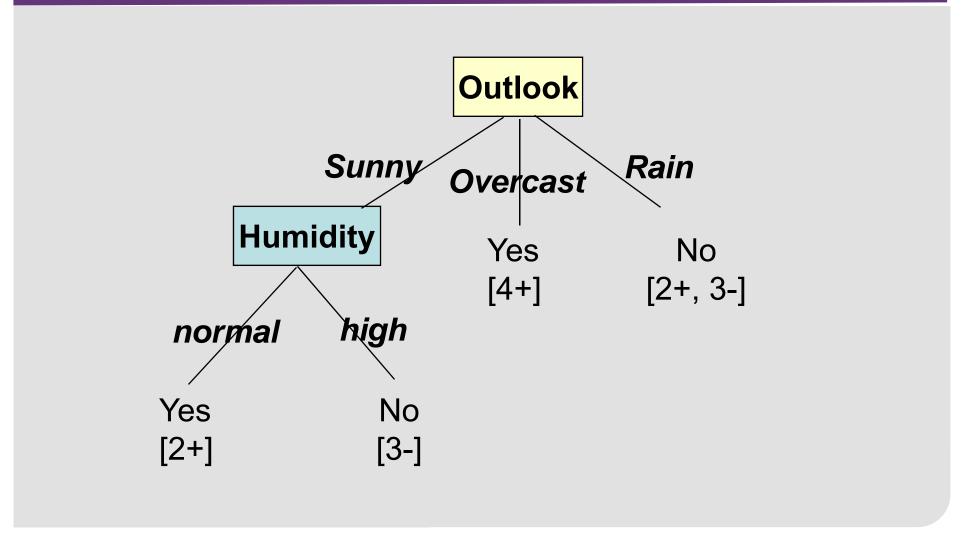
Building a Decision Tree (II)



Day	Temp	Humid	Wind Play	y ball?
d1	Н	Н	weak no	
d2	Н	Н	strong no	
d8	M	Н	weak no	
d9	С	N	weak yes	6
d11	M	N	strong yes	3



Building a Decision Tree (III)





Decision Tree Learning Algorithm (II)

BuildTree(TrainingData)

Split(TrainingData)

```
Split(D)
```

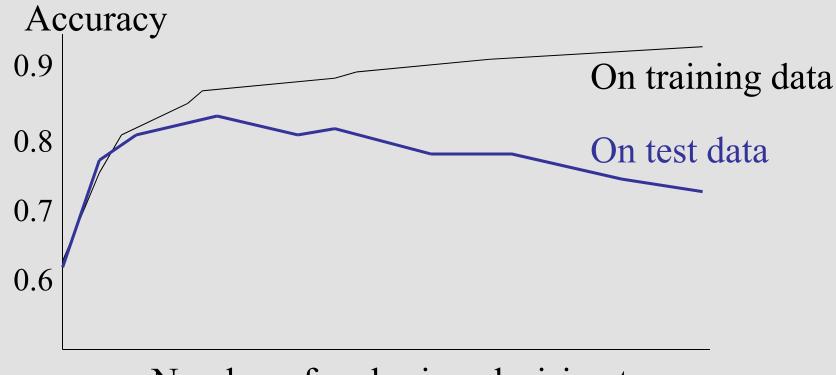
- 1. If all points in D are of the same class Then Return
- 2. For each attribute A
 - a. Evaluate splits on attribute A
- 3. Use best split to partition D into D1, D2, ...
 - a. Split (D1)
 - b. Split (D2) recursive call



Overfitting (I)

The generated tree may *overfit* the training data – try to fit noise or outliers

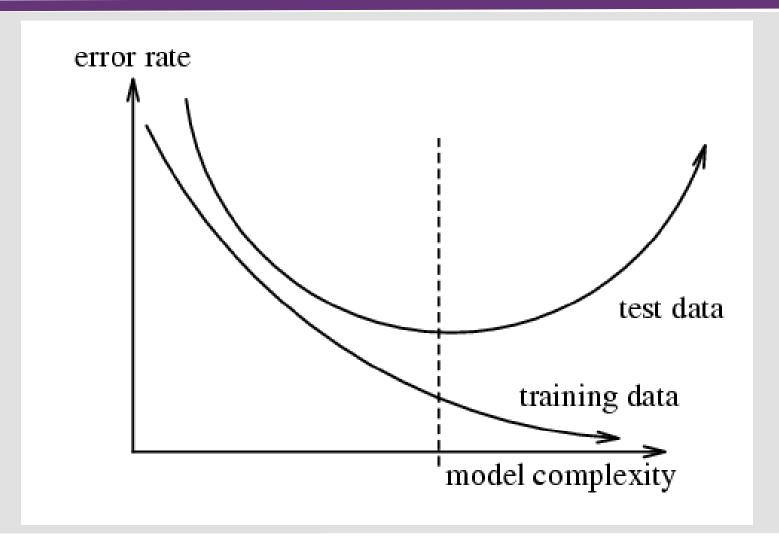
→ reduced accuracy for unseen samples



Number of nodes in a decision tree



Overfitting (II)





Overfitting (III)

- DT is overfitting if there exists another DT' such that
 - DT has a smaller error than DT' on training examples
 - DT has larger error than DT' on test examples
- Causes of overfitting
 - Noisy data
 - Training set is too small

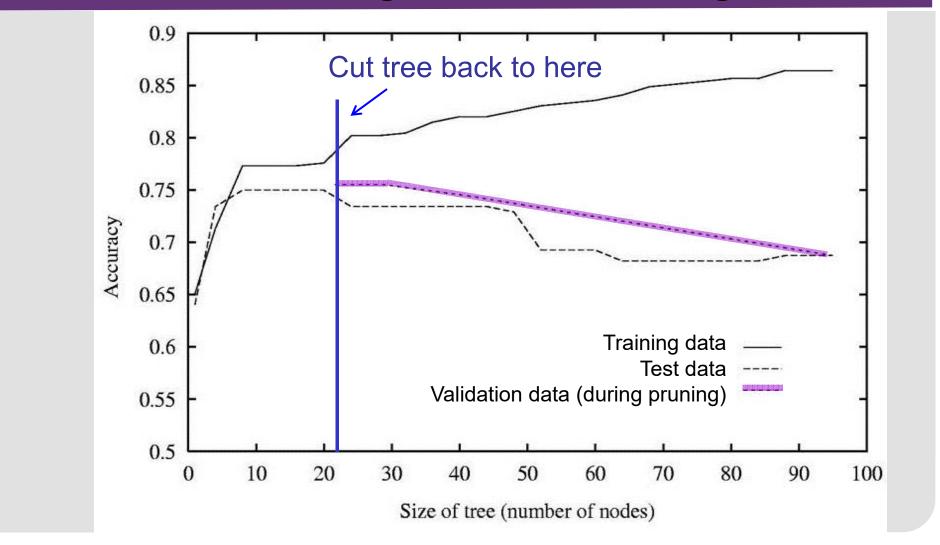


Avoiding Overfitting

- How to prevent overfitting:
 - Pre-pruning: Stop growing the tree if the goodness measure falls below a threshold
 - Post-pruning: Grow the full tree, then prune
 - Regularization: Add complexity penalty to the performance measure
 - > E.g., Complexity = Number of nodes in the tree
- How to select the best tree?
 - Measure performance on training data
 - Measure performance on a separate validation set



Effect of Pruning after Growing the DT





Enhancements to DT Induction

- Handle continuous-valued attributes
 - Input: Use threshold to split
 - Output: Estimate a linear function at each leaf (e.g., mean)
- Handle missing attribute values
 - Categorical
 - > use the most common value of the attribute
 - > probabilistic selection according to the distribution of values
 - Continuous
 - > use the mean of the attribute values
- Reduce overfitting
 - E.g., by post-pruning
- Attribute construction
 - create new attributes based on existing ones that are sparsely represented → reduces sparseness

done manually



Continuous-valued Attributes

Partition the continuous attribute value into intervals

- 1. Fit a distribution to the values for attribute A
 - commonly Normal
- 2. Search for a point to split on
 - perform binary search
 - can split on the same attribute again (lower in the tree)
- 3. Calculate the Information Gain obtained from splitting attribute A at that point



Continuous-valued Attributes – Example

- Humidity has a Normal distribution with mean μ =81.643 and standard deviation σ =10.285
- When we split on 75 (after outlook=sunny), we obtain

normal humidity high humidity $IG(S, h \le 75) = H(S) - \frac{2}{5}H(h \le 75) - \frac{3}{5}H(h > 75)$

$$\frac{10(3, n \le 75) - 11(3) - -11(n \le 75) - -11(n > 75)}{5}$$



Extracting Classification Rules from DTs

- Represent the inferred structure in the form of IF-THEN rules
- One rule is created for each path from the root to a leaf
- Each attribute-value pair along a path forms a conjunction
- The leaf node holds the class prediction



Extracting Rules from DTs – Example

- IF outlook = "sunny" and humidity="high"
 THEN play-ball = "NO"
- IF outlook = "sunny" and humidity="normal"
 THEN play-ball = "YES"
- IF outlook = "overcast" THEN play-ball = "YES"
- IF outlook = "rain" and wind="strong" THEN play-ball = "NO"
- IF outlook = "rain" and wind="weak" THEN play-ball = "YES"







FIT5047 –Intelligent Systems

Naïve Bayes Classifier

Naïve Bayes Classifier (I)

- Based on Bayes rule where $Pr(C_i | V_i) = \frac{Pr(V_i | C_i) Pr(C_i)}{Pr(V_i)}$
 - $-C_i$ is the class of item i
 - $V_i = \{v_{il}, ..., v_{in}\}$ are the values of a set of attributes for item i
 - $-v_{ij}$ is the value of attribute j for item i

$$\Pr(C_i = c \mid v_{i1}, ..., v_{in}) = \frac{\Pr(v_{i1}, ..., v_{in} \mid C_i = c) \Pr(C_i = c)}{\Pr(v_{i1}, ..., v_{in})}$$

Assumes conditional independence of the attribute values for different classes

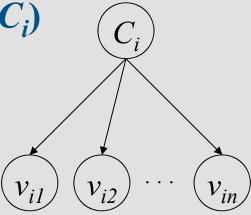
$$\Pr(C_i = c \mid v_{i1}, ..., v_{in}) = \alpha \prod_{k=1}^n \Pr(v_{ik} \mid C_i = c) \Pr(C_i = c)$$

– where α is a normalizing constant



Naïve Bayes Classifier (II)

• The parameters: $Pr(C_i)$ and $Pr(V_i | C_i)$



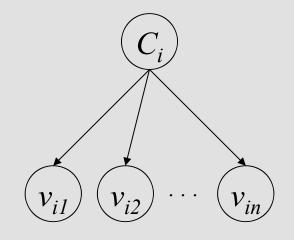
- BIG questions:
 - Prediction/classification: Given that we know the parameters, how do we predict or classify an instance?
 - Learning: How do we learn the parameters?



Naïve Bayes Classifier – Usage

• The predicted class for a new $x_i = (v_{i1}, ..., v_{in})$ is:

$$c^* = \arg\max_{c} \Pr(c \mid v_{i1}, \dots, v_{in})$$
$$= \arg\max_{c} \prod_{k=1}^{n} \Pr(v_{ik} \mid c) \Pr(c)$$



Naïve Bayes Classifier – Example

- 4 attributes outlook, temperature, humidity, wind
- 2 target classes YES/NO (play ball)
- Probability of a class:

<i>c</i> =	YES	NO
$Pr(C_i=c)$	9/14	5/14

obtained from the data

• Calculating $Pr(C_i=YES|v_{i1}=sunny,v_{i2}=hot,v_{i3}=high,v_{i4}=weak)$

$$Pr(C_{i} = YES \mid v_{i1} = sunny, v_{i2} = hot, v_{i3} = high, v_{i4} = weak)$$

$$= \alpha Pr(v_{i1} = sunny \mid C_{i} = YES) \times Pr(v_{i2} = hot \mid C_{i} = YES) \times$$

$$Pr(v_{i3} = high \mid C_{i} = YES) \times Pr(v_{i4} = weak \mid C_{i} = YES) \times Pr(C_{i} = YES)$$

$$= \alpha \frac{2}{9} \times \frac{2}{9} \times \frac{3}{9} \times \frac{6}{9} \times \frac{9}{14} = \alpha \times 0.007$$

This calculation is repeated for all values of c



Learning the Parameters

Estimating the CPTs Pr(c) and Pr(v|c)

- Empirically: use training data
 - Estimate the probability of a class c
 - For each attribute value w_a , estimate the probability of this value from <u>all instances in a class</u>
- · Elicitation: ask a person
 - Usually need domain experts, and sophisticated ways of eliciting probabilities
 - Trouble calibrating



Estimating Parameters Empirically

Maximum Likelihood Estimator (MLE):

$$Pr_{ML}(var = w_a) = \frac{count(var = w_a)}{\# of \ samples} = \frac{|var = w_a|}{\sum_{i=1}^{m} |var = w_i|}$$

- where m is the number of values for var
- Example:









Assume the classes of interest are {H,T}

MLE over all samples

$$Pr_{ML}(H) = \frac{1}{4}, Pr_{ML}(T) = \frac{3}{4}$$

MLE of an attribute over one class

$$\Pr_{ML}(green|T) = \frac{2}{3}$$
, $\Pr_{ML}(yellow|T) = \frac{1}{3}$, $\Pr_{ML}(red|T) = 0$

Maximum Likelihood Principle

- Suppose you have data D, and a probabilistic model parameterized by Θ
 - Need to learn parameter Θ from data D
- Likelihood: The probability of data given the model
- Maximum likelihood: Choose Θ* which maximizes (the log of) the likelihood function:

$$\Theta^* := \operatorname{argmax}_{\Theta} \operatorname{log} \operatorname{Pr}_{\Theta}(D)$$
 Likelihood

Set the derivative to 0

$$\frac{\partial}{\partial \theta}[\log \Pr_{\theta}(D)] = 0 \Rightarrow \frac{count(var = w_a)}{\# of samples}$$



The Sparse Data Problem – Smoothing

- Not all instances are found in the data set or in a particular class
 - If value w_a is not found in class c, $w_a = 0 \rightarrow \text{MLE for } \Pr(w_a|c) = 0$
 - If variable var is not found in the training set, MLE for $Pr(w_a)$ is undefined (denominator is zero)
- Expected Likelihood Estimator (ELE) $\Pr_{EL}(var = w_a) \cong \frac{|w_a| + \varepsilon}{\sum_{i=1}^{m} \{|w_i| + \varepsilon\}} = \frac{|w_a| + \varepsilon}{\sum_{i=1}^{m} |w_i| + m\varepsilon}$
 - where *m* is the number of values for *var*
 - If a variable is not found in the dataset, ELE is 1/m
 - ELE is conservative
- Use Smoothing when estimating the parameters of Naïve Bayes, i.e., Pr(C) and Pr(v|C)



Expected Likelihood Estimation – Example

$$\Pr_{ML}(c = H) \cong \frac{|H|}{\sum_{i=1}^{2} |C_i|} = \frac{1}{\{1+3\}} = \frac{1}{4} \qquad \Pr_{ML}(C) = \left(\frac{1}{4}, \frac{3}{4}\right)$$

$$\Pr_{EL}(c=H) \cong \frac{|H| + \varepsilon}{\sum_{i=1}^{2} |C_i| + 2\varepsilon} = \frac{1 + \varepsilon}{\{1 + 3\} + 2\varepsilon} = \frac{1 + \varepsilon}{4 + 2\varepsilon}$$

$$\text{for } \varepsilon = 1 \qquad \Pr_{EL}(C) = \left(\frac{1}{3}, \frac{2}{3}\right)$$

Enhancements to Naïve Bayes

- Handle continuous-valued attributes
- Handle missing attribute values
 - Categorical
 - > ignore missing values
 - > consider a missing value (?) to be an additional value
 - Continuous
 - > use the mean of the attribute values



Continuous-valued Attributes

- 1. Fit a distribution to the values for attribute A
 - commonly a Normal distribution
- 2. Calculate the probability of the value x of attribute A for item i given the value of class C_i
 - use a Normal density function

$$\Pr(v_{iA} = t \mid C_i = c) \approx \Pr(t - \frac{\varepsilon}{2} \le v_{iA} \le t + \frac{\varepsilon}{2} \mid C_i = c)$$

$$\approx \varepsilon \frac{1}{\sqrt{2\pi\sigma}} e^{-\frac{(t-\mu)^2}{2\sigma^2}}$$



Continuous-valued Attributes – Example

- For class C_i =YES, v_{i2} (temperature) has a Normal distribution with mean μ =73 and standard deviation σ =5.23
- The probability of temperature 85 F given class
 YES is

$$\Pr(v_{i2} = 85 \mid C_i = YES) = \varepsilon \frac{1}{\sqrt{2\pi \times 5.23}} e^{-\frac{(85-73)^2}{2\times 5.23^2}} = 0.0055\varepsilon$$



Decision Trees versus Naïve Bayes

	Decision Trees	Naïve Bayes
Accuracy	depends on the features of the data	
Robustness	good	good
Generality	YES	assumes independent attributes
Model construction speed (scalability)	slower	faster
Model size	bigger	smaller
Interpretability	higher	lower







FIT5047 –Intelligent Systems

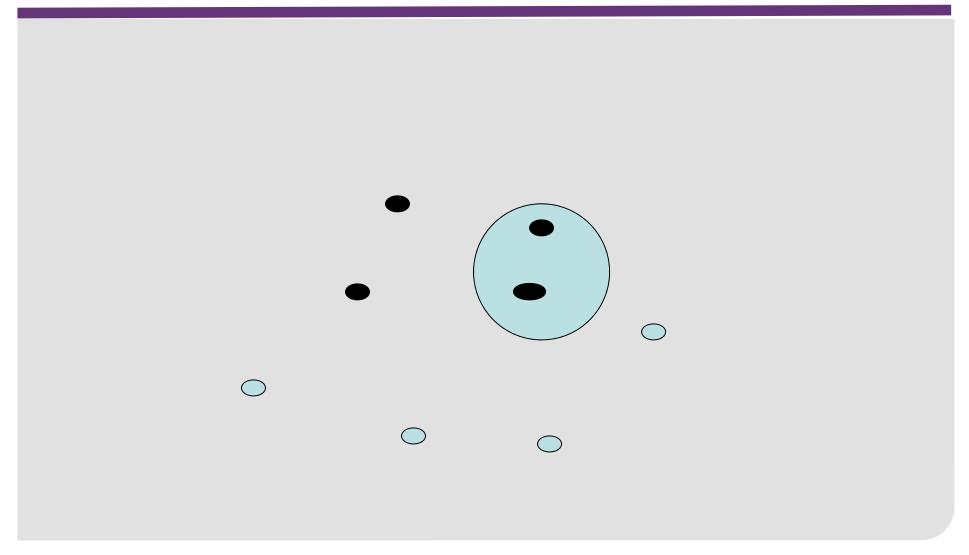
K Nearest Neighbour (k-NN)

k-Nearest Neighbour (I)

- All instances correspond to points in an ndimensional space
- Classification is performed
 - when a new instance arrives
 - by comparing features of the new instance with features of k training instances that are closest to it in the space (nearest neighbours)
- The target function may be discrete or continuous
 - Discrete majority vote of the new instance's neighbours
 - Continuous mean value of the k nearest training examples

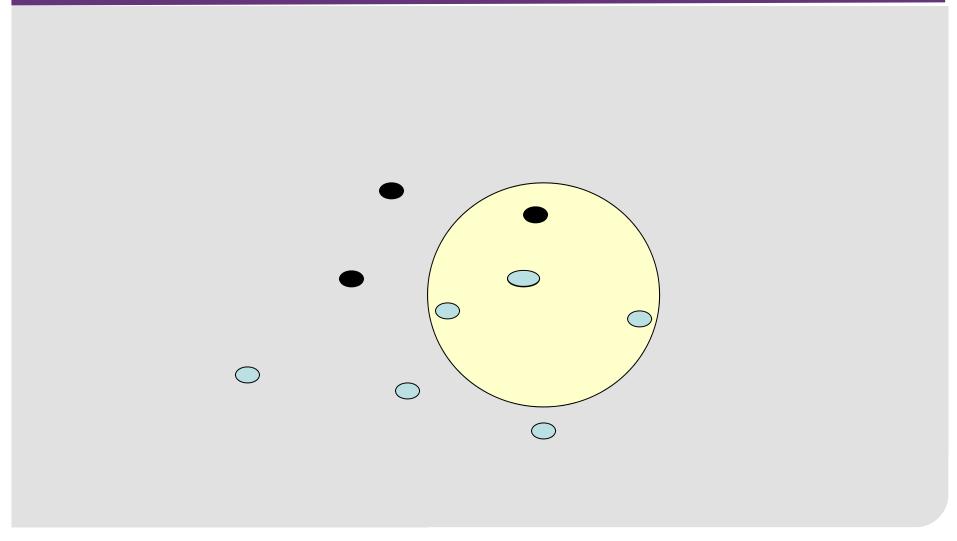


1-Nearest Neighbour





3-Nearest Neighbour





k-Nearest Neighbour (II)

- An instance x_i is represented by $(f_{i1}, f_{i2}, ..., f_{in})$
 - $-f_{i,k}$ is the value of the kth feature for x_i
- Distance measures between two instances x_i and x_i
 - Continuous features: Euclidean distance

$$Ed(x_i, x_j) = \sqrt{\sum_{k=1}^{n} (f_{ik} - f_{jk})^2}$$

Categorical features: Jaccard coefficient

$$Jc(x_i,x_j) = \frac{|\{f_{i1},...,f_{in}\} \cap \{f_{j1},...,f_{jn}\}|}{|\{f_{i1},...,f_{in}\} \cup \{f_{j1},...,f_{jn}\}|}$$

Must be applied to all features

Computing Similarity – Example

Continuous features:

 x_i ={0.7, 30, 80, 10} and x_j ={0.2, 32, 85, 40}

$$Ed(x_i, x_j) = \sqrt{(0.7 - 0.2)^2 + (30 - 32)^2 + (80 - 85)^2 + (10 - 40)^2}$$

Smaller is better!

Categorical features (Jaccard adaptation):

 x_i ={sunny, hot, high, weak} and x_i ={rainy, hot, high, strong}

$$Jc(x_i, x_j) = \frac{|\{hot, high\}|}{|\{sunny, rainy, hot, high, weak, strong\}|} = \frac{2}{6} = 0.33$$

Larger is better!

Need to normalize features



Parameter Learning

- Given the training data and the distance function, there is no training
 - The algorithm memorizes the training data
 - As data comes in, the model size grows
- In some cases, the distance function is parameterized and its parameters are learnt



Other Classification Methods

- Perceptrons
- Support vector machines
- Logistic regression



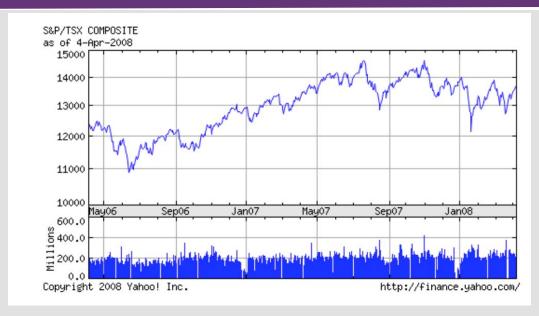




FIT5047 – Intelligent Systems

Regression

Regression



- Given a training set $\{(x_1,t_1),...,(x_m,t_m)\}$
 - $-x_i$ and t_i are continuous \rightarrow regression
 - assume $x_i \in R^n, t_i \in R$
- E.g., t_i is stock price, x_i contains company profit, debt, cash flow, ...



Error Function

- Given a training set $\{(x_1,t_1),...,(x_m,t_m)\}$
 - assume $x_i \in R^n, t_i \in R$



- Linear regression model: $t = w \cdot x + w_{\theta}$
 - we want to learn the parameters $w \in R^n$, $w_0 \in R$
- Error: Square of the difference between the true and predicted target value for x_i

$$E(w) = \sum_{i=1}^{m} (t_i - (w \cdot x_i + w_0))^2$$
truth prediction



Learning the Parameters

Find w* that minimizes the error function

$$\mathbf{w}^* = \operatorname{argmin}_{\mathbf{w}} E(\mathbf{w})$$

- where E(w) is a convex function, so w* is unique
- How to find w*?
 - Set the derivatives to zero: $\frac{\partial}{\partial w_k} E(w) = 0$
 - > For a linear function $t = w_I x + w_0$

$$w_{0} = \frac{\sum t_{i} - w_{1} \sum x_{i}}{N} \qquad w_{1} = \frac{N \sum x_{i} t_{i} - \sum x_{i} \sum t_{i}}{N \sum x_{i}^{2} - (\sum x_{i})^{2}}$$

- Use an iterative algorithm such as gradient descent

Linear Regression – Example (I)

- x_1,t_1 x_2,t_2 x_3,t_3 Training data: {(1,3),(2.1,0.5),(-5,6.2)}
- Linear regression model: $t = w_1 x + w_0$
- Error function:

$$E(w) = (3 - 1w_1 - w_0)^2 + (0.5 - 2.1w_1 - w_0)^2 + (6.2 - (-5)w_1 - w_0)^2$$

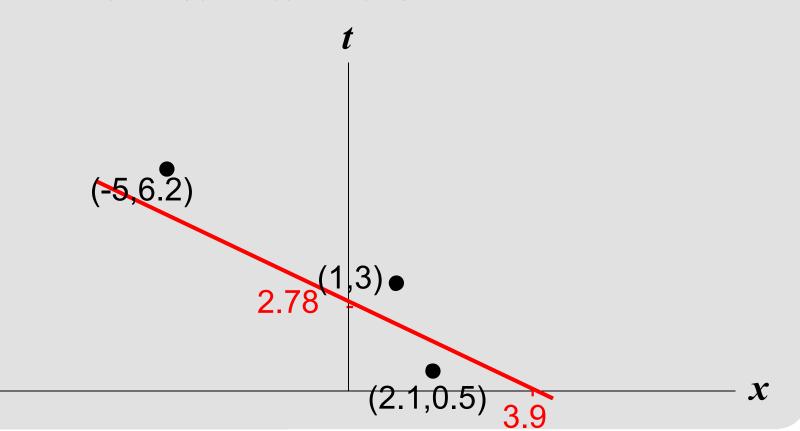
$$w_0 = \frac{(3+0.5+6.2) - w_1(1+2.1-5)}{3} = \frac{9.7 + w_1 1.9}{3} = 2.78$$

$$w_1 = \frac{3(1 \times 3 + 2.1 \times 0.5 + (-5) \times 6.2) - (1 + 2.1 - 5)(3 + 0.5 + 6.2)}{3(1^2 + 2.1^2 + (-5)^2) - (1 + 2.1 - 5)^2} = -0.712$$



Linear Regression – Example (II)

- Training data: {(1,3),(2.1,0.5),(-5,6.2)}
- Function: t = -0.712x + 2.78





Gradient Descent Algorithm

1. initialize w⁰ arbitrarily

- 2. for t = 1,2,... Learning rate
 - a. $\mathbf{w}^t \leftarrow \mathbf{w}^{t-1} \alpha \nabla_{\mathbf{w}} E(\mathbf{w})$

Gradient vector: stack up the partial b. if $|w^t - w^{t-1}| < \varepsilon$ then break derivatives $\frac{\partial E(w)}{\partial w_i}$ in a vector

Illustration of Gradient Descent (I)

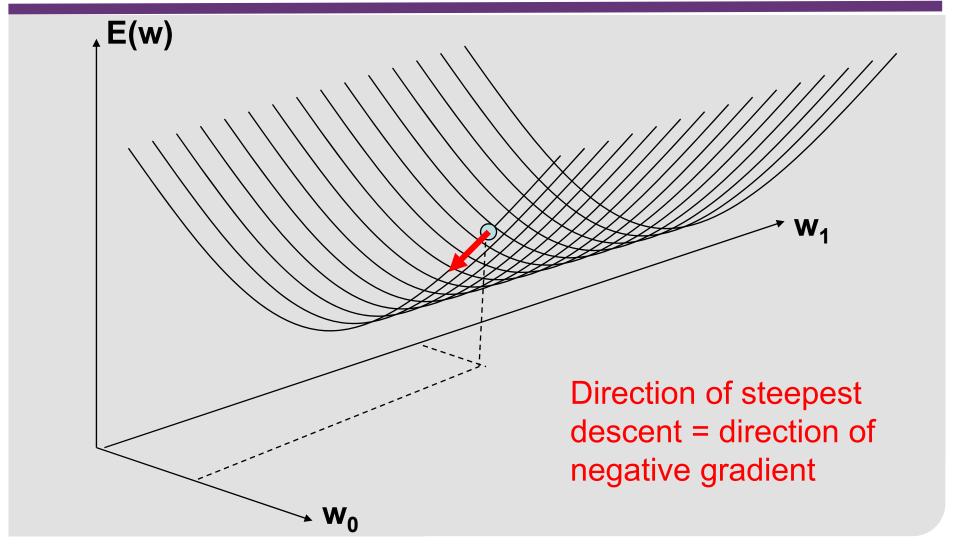
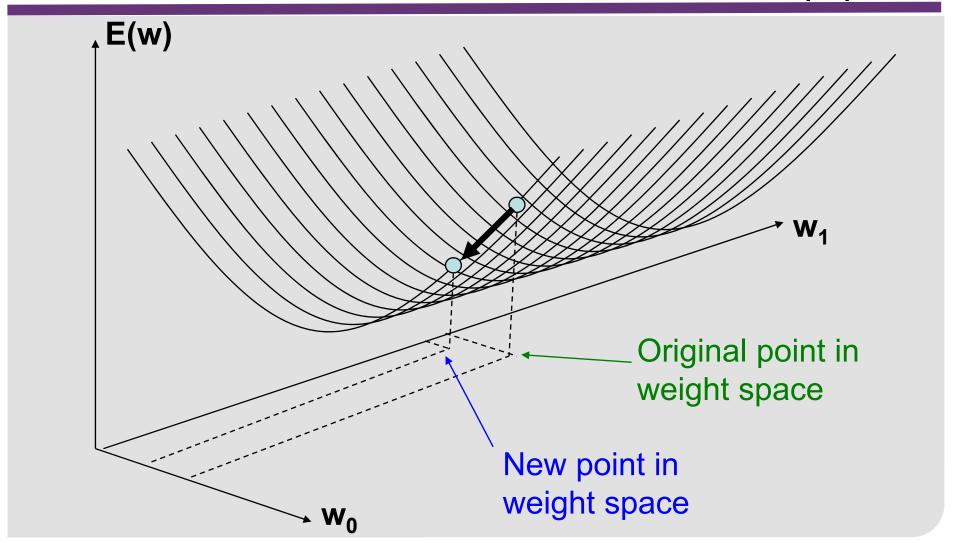




Illustration of Gradient Descent (II)





Linear Regression – Example (II)

- Training data: {(1,3),(2.1,0.5),(-5,6.2)}
- Linear regression model: $t = w_1 x + w_0$
- Error function:

$$E(w) = (3 - 1w_1 - w_0)^2 + (0.5 - 2.1w_1 - w_0)^2 + (6.2 - (-5)w_1 - w_0)^2$$

The gradient:

$$\nabla_{w}E(w) = \begin{bmatrix} \frac{\partial E}{\partial w_{0}} \\ \frac{\partial E}{\partial w_{1}} \end{bmatrix} - 2 \times \begin{bmatrix} (3 - w_{1} - w_{0}) + (0.5 - 2.1w_{1} - w_{0}) + (6.2 + 5w_{1} - w_{0}) \\ (3 - w_{1} - w_{0}) + 2.1(0.5 - 2.1w_{1} - w_{0}) - 5(6.2 + 5w_{1} - w_{0}) \end{bmatrix}$$



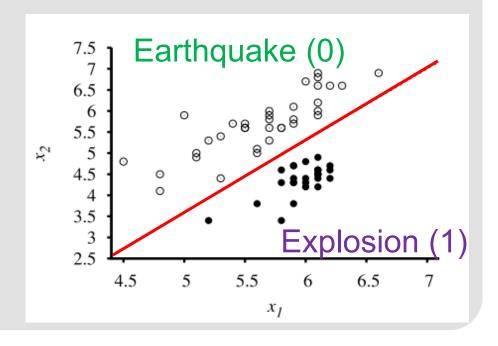
Classification with Linear Regression

- The regressor is a decision boundary between two separable classes
 - A linear decision boundary is called a *linear separator*,
 in this case the data are *linearly separable*
- Example:
 - Linear separator

$$-4.9+1.7x_{1}-x_{2}$$

Classification

$$h(x_1, x_2) = \begin{cases} 1 & \text{if } w \cdot x > 0 \\ 0 & \text{otherwise} \end{cases}$$





Summary – Supervised Machine Learning

- Various techniques for supervised machine learning
 - where we are given labeled training data
- Discrete:
 - Decision trees (non-parametric)
 - Naïve-Bayes (parametric)
 - K nearest neighbour (non-parametric)
- Continuous:
 - Regression
 - > Models based on least square error (parametric)







FIT5047 – Intelligent Systems

Unsupervised Machine Learning – Clustering

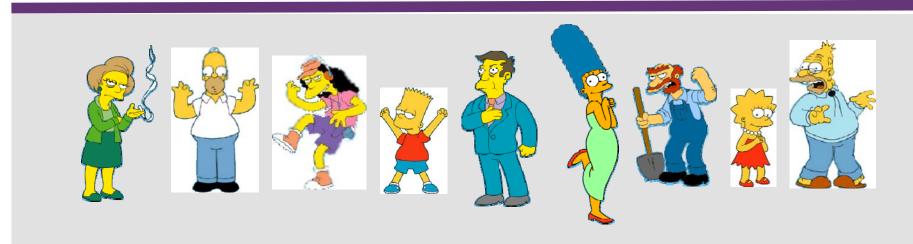
Slides adapted from Eamonn Keogh, Dan Klein and Kevin Korb

What is Clustering?

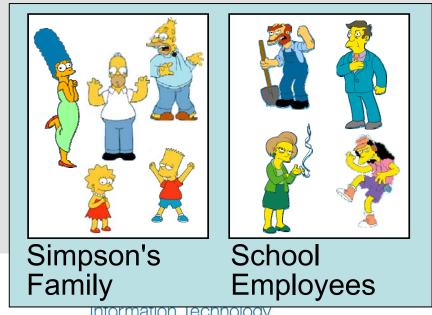
- Organizing data into classes such that there is
 - high intra-class similarity
 - low inter-class similarity
- Finding the class labels and the number of classes directly from the data

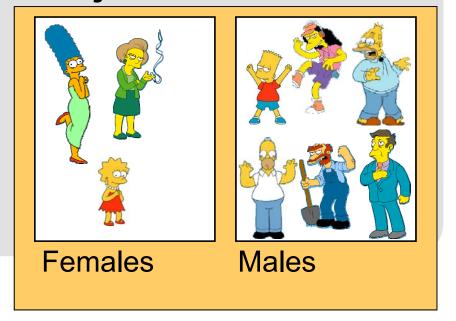


What is a Natural Grouping of these Objects?



Clustering is subjective





Information Technology

What is Similarity?

- The quality or state of being similar; likeness; resemblance; as, a similarity of features.
 Webster's dictionary
- Similarity is hard to define, but...
 "We know it when we see it"





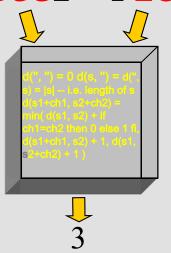
Distance Measures (I)

Definition: Let O_1 and O_2 be two objects from the universe of possible objects. The distance (dissimilarity) between O_1 and O_2 is a real number denoted by $D(O_1, O_2)$ Peter Piotr 0.23 342.7



Distance Measures (II)

Peter Piotr



When we peek inside one of these black boxes, we see a function of two variables

What properties should a distance measure have?

•
$$D(A,B) = D(B,A)$$

• D(A,A) = 0

• D(A,B) = 0 iff A = B

• $D(A,B) \leq D(A,C) + D(B,C)$

Symmetry

Constancy of Self-Similarity

Positivity (Separation)

Triangle Inequality



Measuring Similarity

Edit distance: Transform one of the objects into the other, and measure how much effort it took

Distance between Patty and Selma:

Change dress color 1 point

Change earring shape 1 point

Change hair part 1 point

D(Patty,Selma) = 3

Distance between Marge and Selma:

Change dress color 1 point

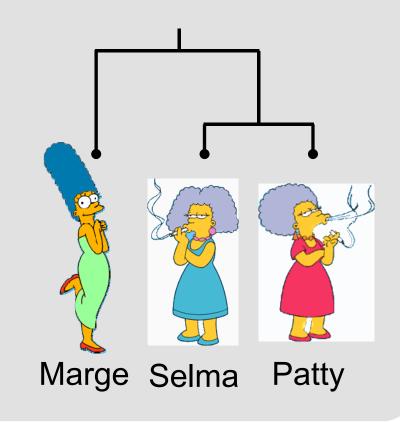
Add earrings 1 point

Decrease height 1 point

Take up smoking 1 point

Gain weight 1 point

D(Marge, Selma) = 5





Two Types of Clustering

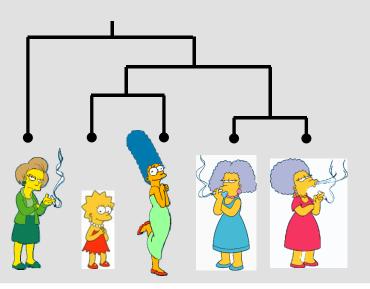
- Partitional: Construct various partitions and then evaluate them by some criterion
- Hierarchical: Create a hierarchical decomposition of the set of objects using some criterion

Partitional





Hierarchical







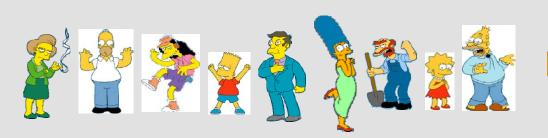


FIT 5047 – Intelligent Systems

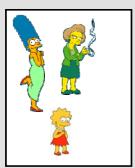
K-means Algorithm

Partitional Clustering

- Non-hierarchical, each instance is placed in exactly one of K non-overlapping clusters
- Produces only one set of clusters
 - → the user normally has to input the desired number of clusters K

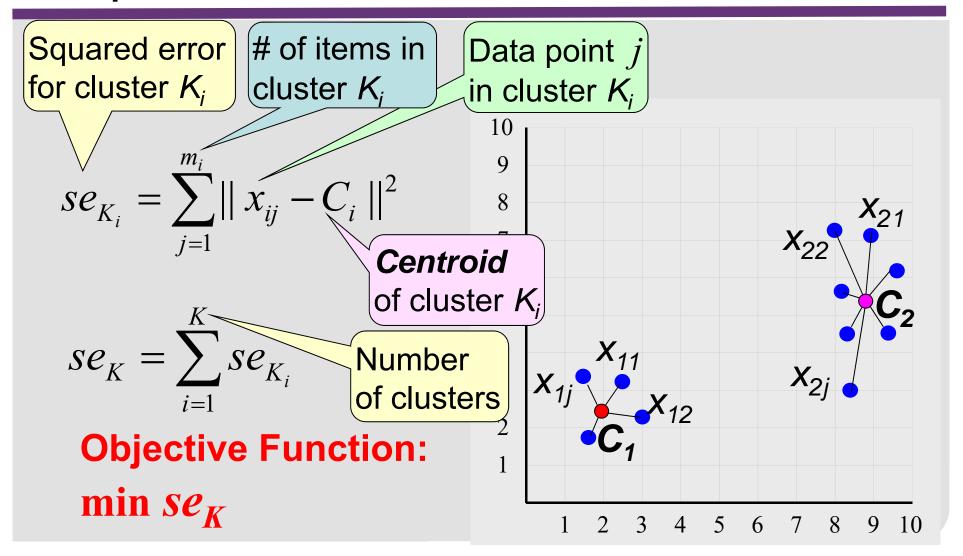








Squared Error





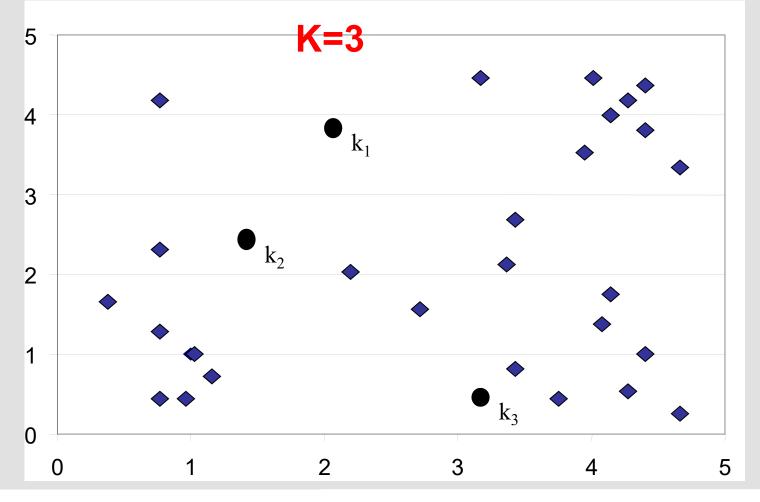
K-means Algorithm

- 1. Decide on a value for k
- 2. Initialize the *k* cluster *centroids* (randomly, if necessary)
- 3. Decide the class memberships of the *N* objects by assigning them to the nearest cluster centroid
- 4. Re-estimate the *k* cluster centroids, assuming the memberships found above are correct
- 5. If none of the *N* objects changed membership in the last iteration, exit
- 6. (Otherwise) goto Step 3



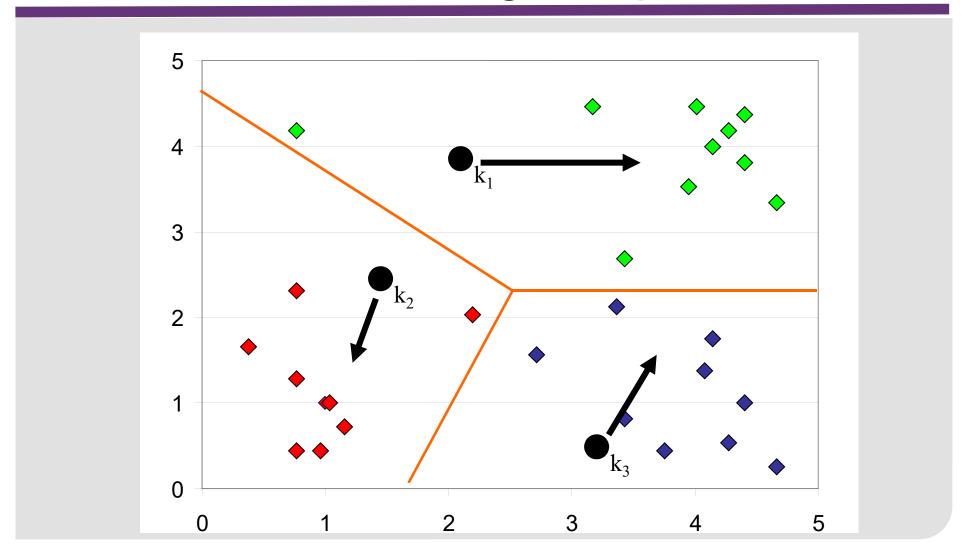
K-means Clustering: Steps 1 and 2

Distance Metric: Euclidean Distance



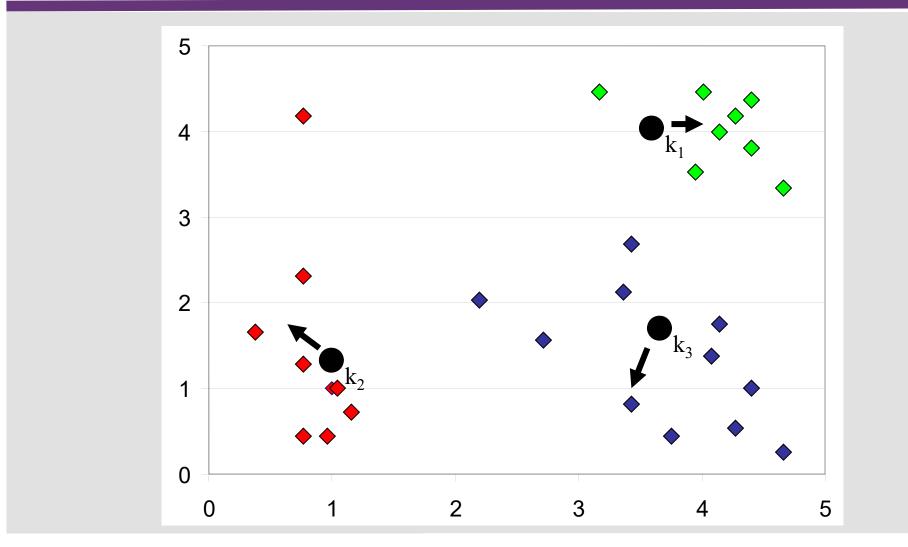


K-means Clustering: Steps 3 and 4



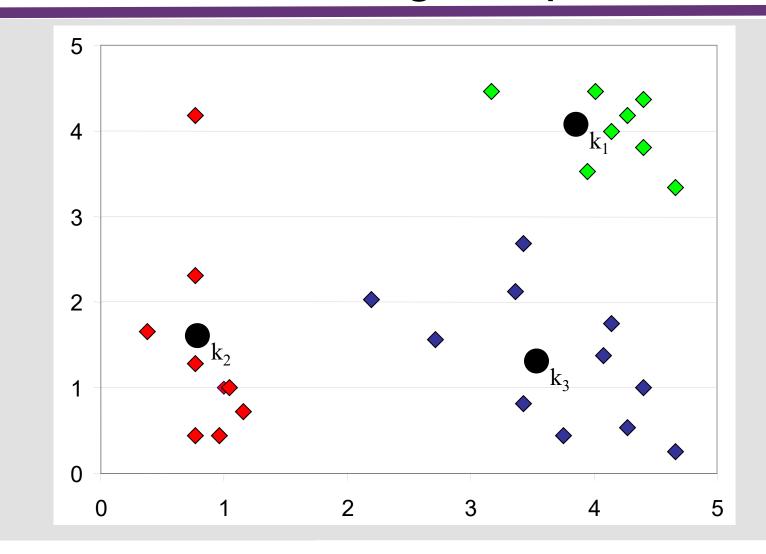


K-means Clustering: Steps 3 and 4





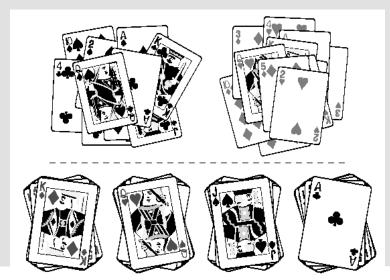
K-means Clustering: Step 5





The Meaning of K

- Clusters describe underlying structure in data
 - but structures in data can exist at different levels
- In many cases, there is no a priori reason to select a particular value for k
 - Should consider several k-s, but different values of k can lead to different clusterings



$$k = 2$$

$$k = 4$$



Normalization

- Problem: distance will be dominated be attributes with large magnitude
- Example: in which cluster do we put age=25 & income=\$30,000?
 - Centroid 1: age=26 & income=\$25,000
 - Centroid 2: age=80 & income=\$34,500
- Solution: normalize the data



Summary of the K-Means Algorithm

Advantages

- Relatively efficient: O(tkn), where n is # objects, k is
 # clusters, and t is # iterations. Normally, k, t << n
- Global optimum may be found using techniques such as deterministic annealing and genetic algorithms

Disadvantages

- Need to specify k in advance
- Applicable only when mean is defined
 - > What about categorical data?
- Does not deal well with overlapping clusters
- Unable to handle noisy data and outliers
 - > Outliers can pull cluster centers



WEKA

www.cs.waikato.ac.nz/ml/weka

Several classifiers

- weka.classifiers
- bayes.NaiveBayes Naïve Bayes
- trees.DecisionStump decision trees with one split only
- trees.J48 decision trees
- lazy.IBk k nearest neighbour
- Several clustering algorithms
 - K-means



Reading

- Supervised learning Russell, S. and Norvig, P. (2003), Artificial Intelligence – A Modern Approach (3rd ed), Prentice Hall
 - Decision trees Chapter 18.1-18.2
 - Naïve Bayes Chapter 20.1-20.2
 - Regression Chapter 18.6.1
 - k-NN– Chapter 18.8.1
- Naïve Bayes
 - John, G. and Langley, P. (1995), Estimating continuous distributions in Bayesian Classifiers. In *Proceedings of the 11th* Conference on Uncertainty in Artificial Intelligence, pp. 338-345
- Unsupervised learning and K-means clustering
 - http://en.wikipedia.org/wiki/K-means clustering
 - https://sites.google.com/site/dataclusteringalgorithms/k-meansclustering-algorithm

